

# 第五讲 PyTorch Tutorial

王文中

安徽大学计算机学院

# Deep Learning with PyTorch

Eli Stevens  
Luca Antiga  
Thomas Viehmann  
Foreword by Soumith Chintala

 MANNING

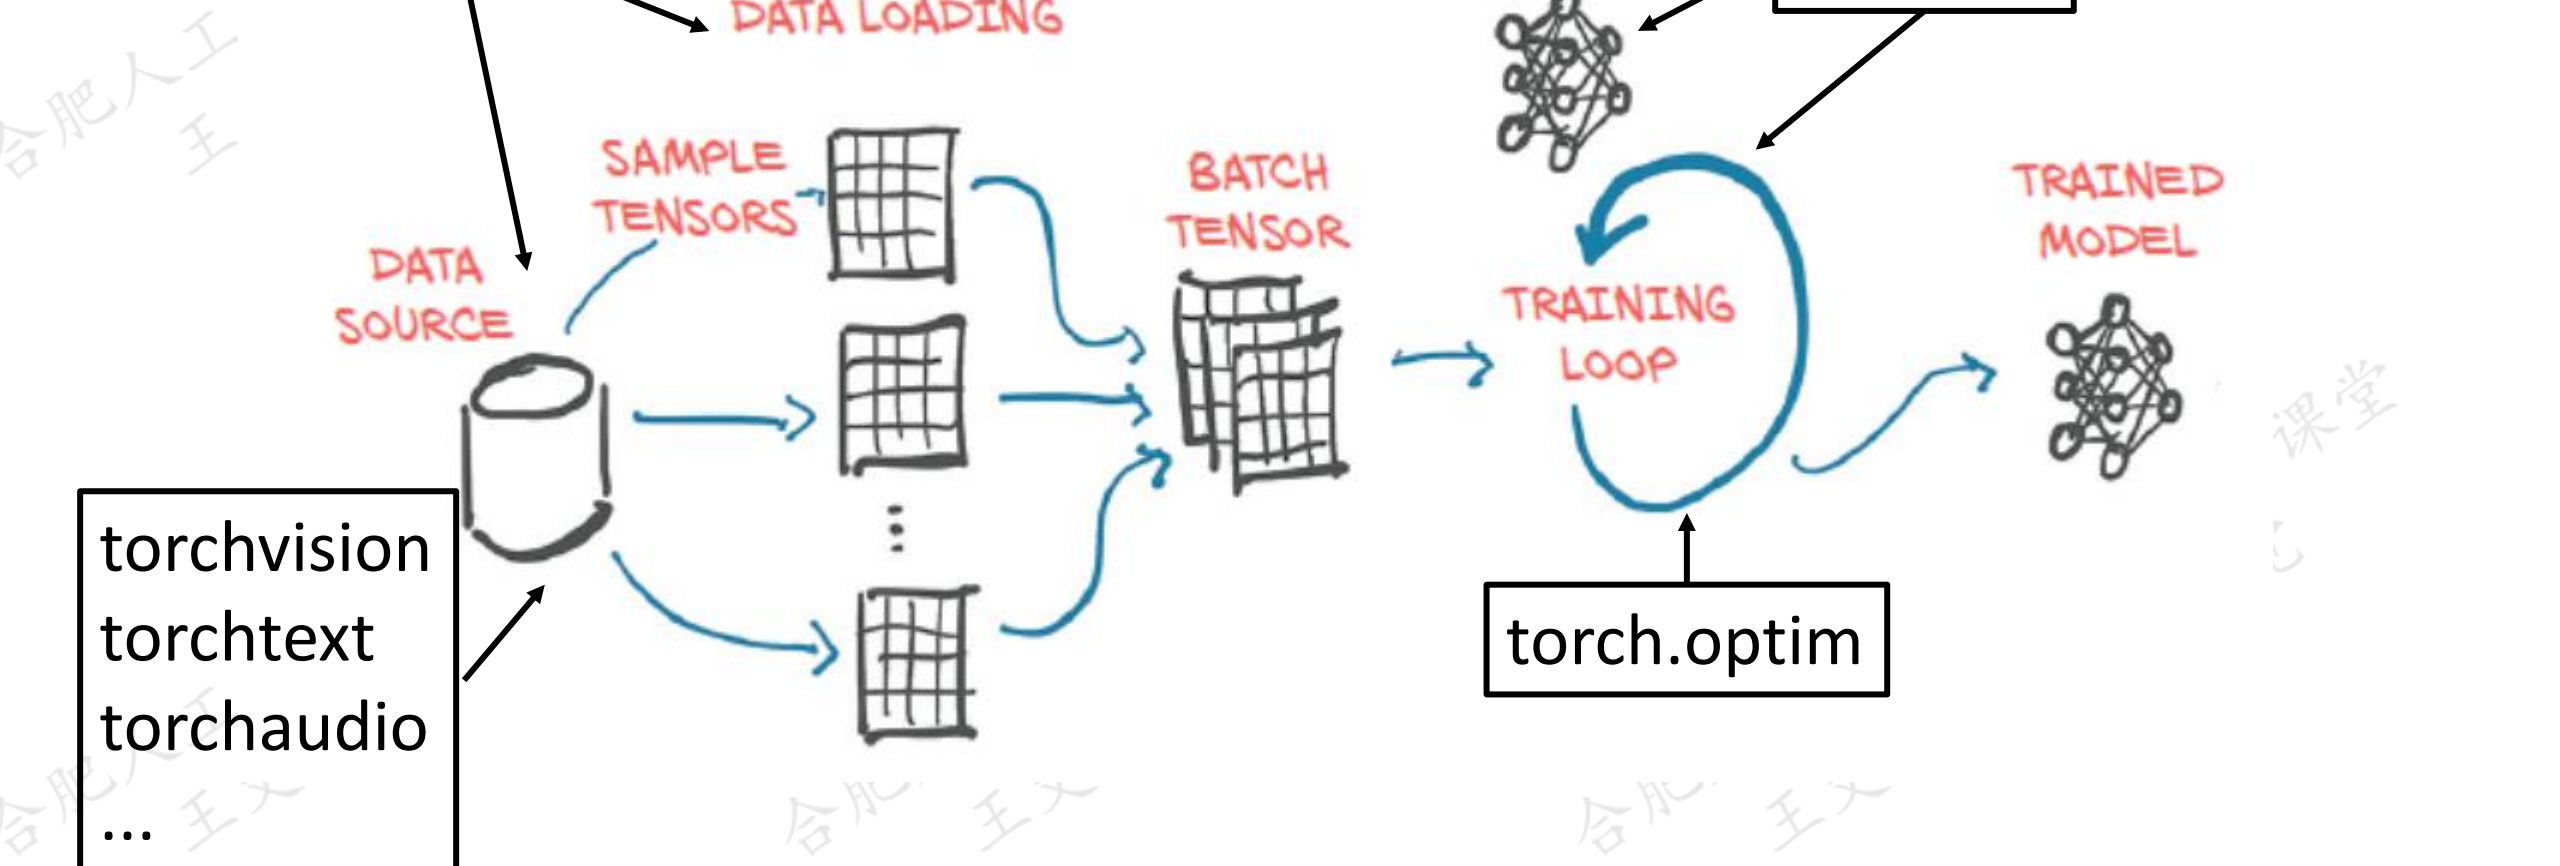


<https://www.manning.com/books/deep-learning-with-pytorch>

# 大纲



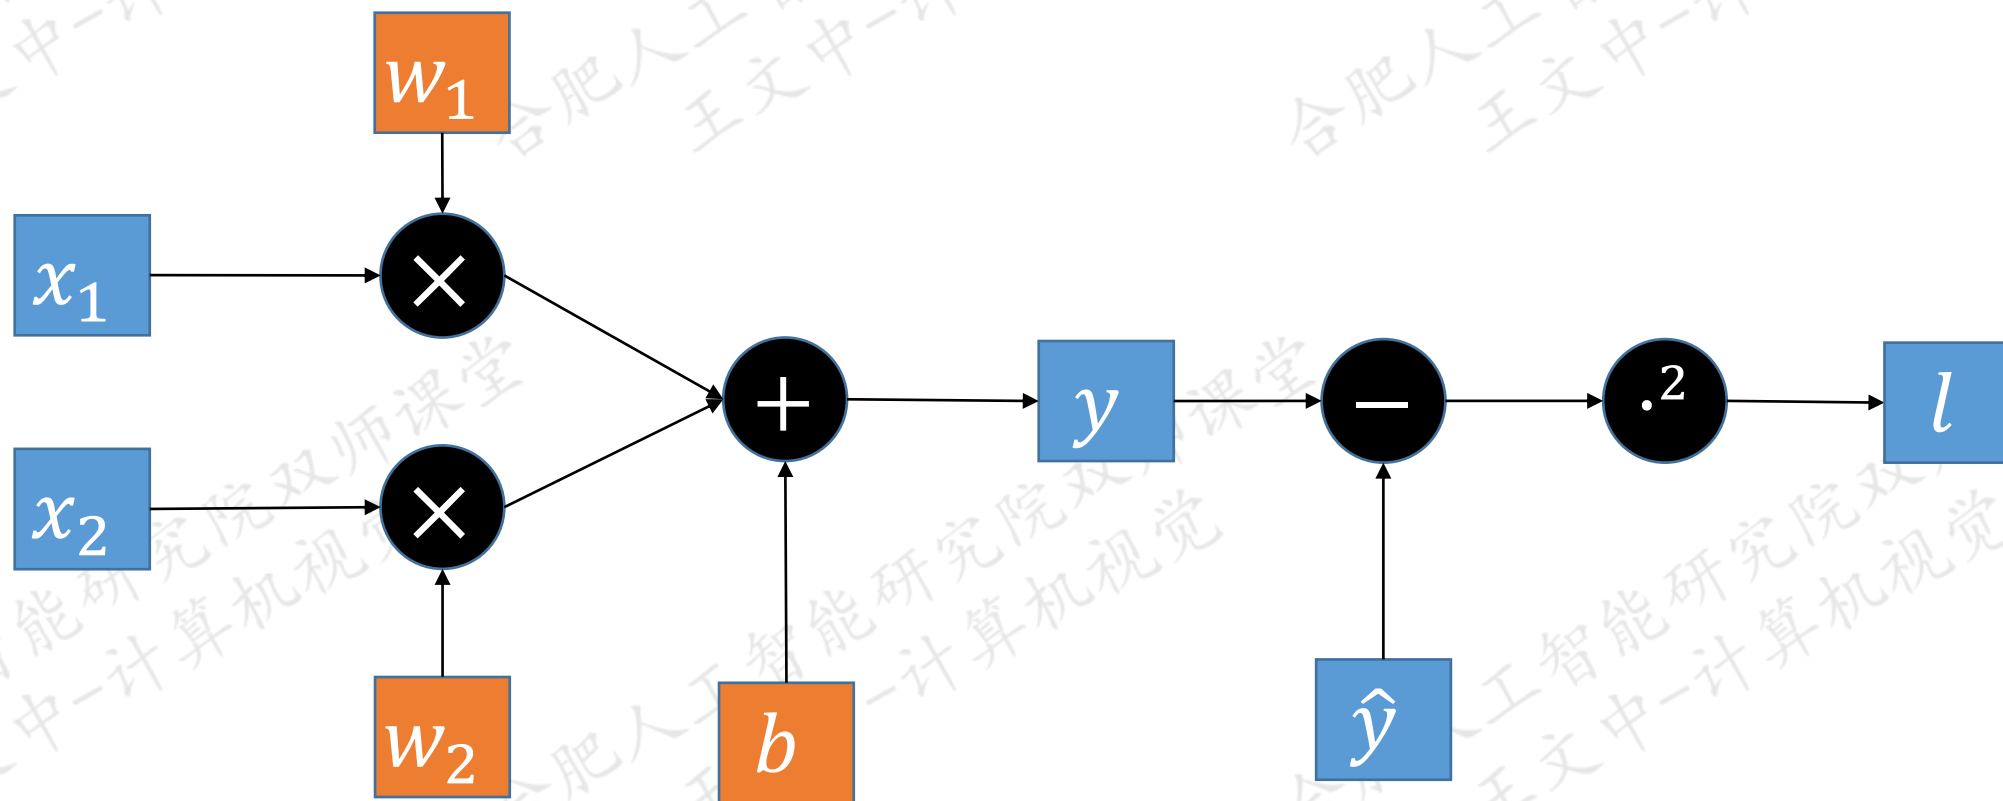
- 核心数据结构Tensor
- 多层感知机编程
- 卷积神经网络编程



# 计算图



$$y = h(x; w, b) = w_1 \times x_1 + w_2 \times x_2 + b$$
$$l = (y - \hat{y})^2$$



# 用张量表示数据与参数



$$y = h(x; w, b) = w_1 \times x_1 + w_2 \times x_2 + b$$
$$l = (y - \hat{y})^2$$

$$x = \begin{pmatrix} x_1 \\ x_2 \end{pmatrix}, w = \begin{pmatrix} w_1 \\ w_2 \end{pmatrix}, b$$

$$y = w^T x + b$$

$$l = (y - \hat{y})^2$$

$$X = (x^{(1)}, x^{(2)}, \dots, x^{(n)})$$

$$\hat{Y} = (y^{(1)}, y^{(2)}, \dots, y^{(n)})$$


$$Y = w^T X + b$$

$$L = (Y - \hat{Y})^2$$

# 计算:张量流(Tensor Flow)



$$y = h(x; w, b) = w_1 \times x_1 + w_2 \times x_2 + b$$
$$l = (y - \hat{y})^2$$





# 1.核心数据结构:Tensor



# Tensor: 多维数组



3

SCALAR

0D

$$\begin{bmatrix} 4 \\ 1 \\ 5 \end{bmatrix}$$

VECTOR

$X[2] = 5$

1D

$$\begin{bmatrix} 4 & 6 & 7 \\ 7 & 3 & 9 \\ 1 & 2 & 5 \end{bmatrix}$$

MATRIX

$X[1, 0] = 7$


2D

$$\begin{bmatrix} 5 & 7 & 1 \\ 9 & 4 & 3 \\ 3 & 5 & 2 \end{bmatrix}$$

TENSOR

$X[0, 2, 1] = 5$

3D



TENSOR


$X[1, 3, \dots, 2] = 4$

|

N-D DATA  $\rightarrow$  N INDICES

# 1.1 创建与操作Tensor

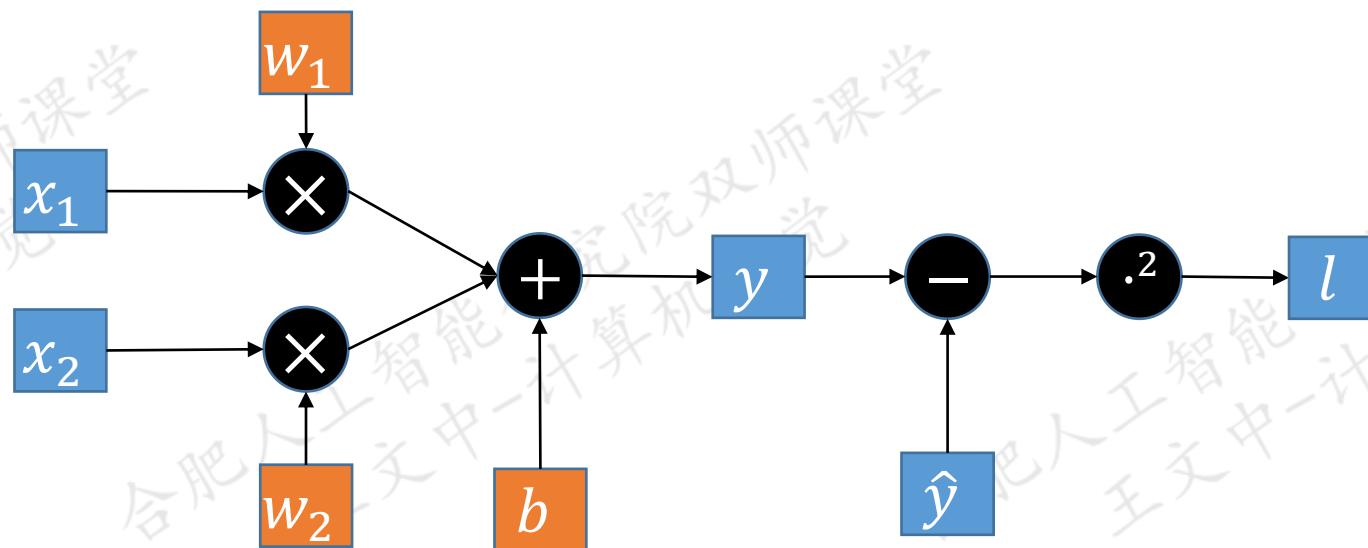
# Tensor的属性



# 创建一个Tensor



- 1.使用torch的创建算子(Creation Ops)
- 2.通过对Tensor的运算创建



# 创建一个Tensor



```
a = torch.tensor(1.0)
print(a)
print(a.dim())
print(a.shape)#print(a.size())
print(a.dtype)
print(a.item())
```

```
tensor(1.)
0
torch.Size([1])
torch.float32
1.0
```

```
b = torch.tensor([1,2,3])
print(b)
print(b.dim())
print(b.shape)
print(b.dtype)
print(b.numpy())
```

```
tensor([1, 2, 3])
1
torch.Size([3])
torch.int64
[1 2 3]
```

```
x = np.array([0,1])
c = torch.from_numpy(x)
```

```
tensor([0, 1], dtype=torch.int32)
```

# 创建一个Tensor



```
c = torch.tensor([[1,2],[3,4],[5,6]],dtype = torch.int16)
print(c)
print(c.dim())
print(c.shape)
print(c.dtype)
print(c.numpy())
```

```
tensor([[1, 2],
        [3, 4],
        [5, 6]], dtype=torch.int16)
2
torch.Size([3, 2])
torch.int16
[[1 2]
 [3 4]
 [5 6]]
```

# 创建一个Tensor



```
d = torch.zeros((3,4))
```

```
tensor([ [0., 0., 0., 0.],  
        [0., 0., 0., 0.],  
        [0., 0., 0., 0.]])
```

```
e = torch.randn((2,2,3,4))
```

```
tensor([[[[-1.0510,  0.3902,  0.1810,  1.0036],  
          [-1.4556,  0.4581, -0.5105, -1.3028],  
          [-0.8999,  0.3423,  0.0936,  1.3038]],  
        [[-0.7691,  0.4797, -1.8662, -0.7550],  
          [ 0.0411,  0.5273, -0.1488,  1.0463],  
          [ 1.1466, -0.3363,  0.7472, -1.2813]]],  
       [[[[-0.3448,  0.3213,  0.2527,  1.2449],  
          [ 0.7206, -1.5667,  0.2489, -1.9135],  
          [ 0.0946, -0.1390,  0.1650,  1.1801]],  
        [[-0.9789,  0.7678, -0.2877, -1.1634],  
          [ 0.7762,  0.7401,  0.2081,  0.3059],  
          [ 0.5780, -0.2145,  0.8056,  0.8838]]]])
```

```
e.stride() (24, 12, 4, 1)
```



# 设置Tensor的设备属性



```
a = torch.tensor([1,2],device = 'cuda')  
b = torch.tensor([1,2]).to(device='cuda')
```

# 访问Tensor的元素



```
a = torch.tensor(1.0)
a[0]
```

**IndexError** Traceback (most recent call last)

```
<ipython-input-68-292e29054dae> in <module>()
      1 a = torch.tensor(1.0)
----> 2 a[0]
```

**IndexError:** invalid index of a 0-dim tensor. Use tensor.item() to convert a 0-dim tensor to a Python number

```
a.item()
type(a.item())
```

```
1.0
float
```

```
a.numpy()
```

```
array(1., dtype=float32)
```

```
a = torch.tensor([1.0,2.0])
a.item()
```

```
a.numpy()
```

```
array([1., 2.], dtype=float32)
```

**ValueError** Traceback (most recent call last)

```
<ipython-input-121-cd1bda83583f> in <module>()
----> 1 a.item()
```

**ValueError:** only one element tensors can be converted to Python scalars

# 访问Tensor的元素



```
b = torch.tensor([[[1,2],[3,4]],  
                  [[5,6],[7,8]],  
                  [[9,10],[11,12]]  
                  ],dtype = torch.int32)
```

b.shape → torch.Size([3, 2, 2])

b.storage() →


b[0] → tensor([[1, 2],  
 [3, 4]])

b[:1] → tensor([[[1, 2],  
 [3, 4]]])

b[:,0,0] → tensor([1, 5, 9])

1  
2  
3  
4  
5  
6  
7  
8  
9  
10  
11  
12  
[torch.LongStorage of size 12]

# 访问Tensor的元素



# 访问Tensor的元素



```
a = torch.ones((3))  
b = a[:2]  
print(a)  
print(b)
```

```
tensor([1., 1., 1.])  
tensor([1., 1.])
```

```
b[0] = 2  
print(a)  
print(b)
```

```
tensor([2., 1., 1.])  
tensor([2., 1.])
```

```
c = a[:2].clone()  
c[0] = 3  
print(a)  
print(c)
```

```
tensor([2., 1., 1.])  
tensor([3., 1.])
```

# 访问Tensor的元素:掩膜



```
X = torch.randn((2,3))
mask = torch.randn((2,3))>0
X1 = X[mask]
X2 = torch.masked_select(X,mask)
print('X=',X.numpy())
print('mask=',mask.numpy())
print('X1=',X1.numpy())
print('X2=',X2.numpy())
```

```
X= [[-1.7936143  -1.0867784   0.6989719 ]
     [ 0.32036787 -0.29389462  0.6172063 ]]
```

```
mask= [[False  True  True]
        [False  True  True]]
```

```
X1= [-1.0867784   0.6989719  -0.29389462  0.6172063 ]
```

```
X2= [-1.0867784   0.6989719  -0.29389462  0.6172063 ]
```

# 访问Tensor的元素:掩膜



```
col_mask = torch.randn((3))>0
print('col_mask=',col_mask.numpy())
X3 = X[:,col_mask]
X4 = torch.masked_select(X,col_mask)
print('X3=',X3.numpy())
print('X4=',X4.numpy())
```

```
X= [[-1.7936143  -1.0867784   0.6989719 ]
     [ 0.32036787 -0.29389462  0.6172063 ]]
```

```
col_mask= [False  True  True]
```

```
X3= [[-1.0867784   0.6989719 ]
      [-0.29389462  0.6172063 ]]
```

```
X4= [-1.0867784   0.6989719  -0.29389462  0.6172063 ]
```



# 访问Tensor的元素:条件选择



```
x = torch.randn(1, 3)
y = torch.randn(1, 3)
z = torch.where(x>y,x,y)
```

```
x tensor([[ 0.0684, -1.5483,  2.1476]])
y tensor([[ -1.0761, -0.0047, -2.0822]])
z tensor([[ 0.0684, -0.0047,  2.1476]])
```

# Transpose



```
a = torch.tensor([[[[1,2],[3,4]],  
                  [[5,6],[7,8]],  
                  [[9,10],[11,12]]  
                ]])  
b = a.transpose(0,1)  
print(a.stride()) (4, 2, 1)  
print(b.stride()) (2, 4, 1)
```

```
tensor([[[[ 1,  2],  
          [ 5,  6],  
          [ 9, 10]],  
        [[[ 3,  4],  
          [ 7,  8],  
          [11, 12]]]])
```

1	2	3	4	5	6	7	8	9	10	11	12
---	---	---	---	---	---	---	---	---	----	----	----

# Transpose



```
a = torch.tensor([[[1,2],[3,4]],  
                  [[5,6],[7,8]],  
                  [[9,10],[11,12]]  
                  ])  
b = a.transpose(0,2)
```

```
tensor([[[[1,2],[3,4]],  
         [[5,6],[7,8]],  
         [[9,10],[11,12]]]])
```



```
tensor([[[[ 1,  5,  9],  
          [ 3,  7, 11]],  
        [[ 2,  6, 10],  
          [ 4,  8, 12]]]])
```

# Transpose



```
a = torch.tensor([[3,1,2],[4,1,7]])
```

```
b = a.transpose(0,1)
```

```
print(a[0,0])
```

```
tensor(3)
```

```
b[0,0] = -1
```

```
print(a[0,0])
```

```
tensor(-1)
```

```
a = torch.tensor([[3,1,2],[4,1,7]])
```

```
b = a.clone().transpose(0,1)
```

```
print(a[0,0])
```

```
tensor(3)
```

```
b[0,0] = -1
```

```
print(a[0,0])
```

```
tensor(3)
```

# Squeeze/unsqueeze



```
a = torch.tensor([[[[1,2,3]]]])
```

```
b = a.squeeze(dim=0)
```

```
c = a.squeeze()
```

```
d = b.unsqueeze(dim=1)
```

a.size()	torch.Size([1, 1, 3])
----------	-----------------------

b.size()	torch.Size([1, 3])
----------	--------------------

c.size()	torch.Size([3])
----------	-----------------

d.size()	torch.Size([3, 1])
----------	--------------------

```
x = torch.randn(2, 3)
x1 = torch.cat((x,x),dim = 0)
x2 = torch.cat((x,x),dim = 1)
```

```
x
tensor([[ -0.4173,  0.0633, -1.3320],
        [ 0.5852, -2.0193,  0.2838]])
```

```
x1
tensor([[ -0.4173,  0.0633, -1.3320],
        [ 0.5852, -2.0193,  0.2838],
        [ -0.4173,  0.0633, -1.3320],
        [ 0.5852, -2.0193,  0.2838]])
```

```
x2
tensor([[ -0.4173,  0.0633, -1.3320, -0.4173,  0.0633, -1.3320],
        [ 0.5852, -2.0193,  0.2838,  0.5852, -2.0193,  0.2838]])
```

# reshape



```
T = torch.arange(0,12)
V1 = T.reshape((2,6))#也可以用torch.reshape(T,(2,6))
V2 = T.reshape((3,4))
print('T=',T)
print('V1=',V1)
print('V2=',V2)
```

```
T= tensor([ 0,  1,  2,  3,  4,  5,  6,  7,  8,  9, 10, 11])

V1= tensor([[ 0,  1,  2,  3,  4,  5],
            [ 6,  7,  8,  9, 10, 11]])

V2= tensor([[ 0,  1,  2,  3],
            [ 4,  5,  6,  7],
            [ 8,  9, 10, 11]])
```



# permute



```
A = torch.arange(24).reshape((2,3,4))
print(A.shape)
print('A=',A)

B = A.permute((2,1,0))
print(B.shape)
print('B=',B)

A[0,0,1] = -5
print('B=',B)#注意B[1,0,0]的值变为-5
```

```
torch.Size([2, 3, 4])
A= tensor([[[ 0,  1,  2,  3],
            [ 4,  5,  6,  7],
            [ 8,  9, 10, 11]],
          [[12, 13, 14, 15],
            [16, 17, 18, 19],
            [20, 21, 22, 23]]])
torch.Size([4, 3, 2])
B= tensor([[[ 0, 12],
            [ 4, 16],
            [ 8, 20]],
          [[ 1, 13],
            [ 5, 17],
            [ 9, 21]],
          [[ 2, 14],
            [ 6, 18],
            [10, 22]],
          [[ 3, 15],
            [ 7, 19],
            [11, 23]]]])
```

# 张量运算



```
A = torch.tensor([1,2,3],dtype = torch.float32)
B = torch.tensor([-1,1,-1],dtype = torch.float32)
```

```
C = A + B
```

```
D = A - B
```

```
E = A*B
```

```
F = A/B
```

```
A= [1.  2.  3.]
```

```
B= [-1.   1.  -1.]
```

```
A+B= [0.  3.  2.]
```

```
A-B= [2.  1.  4.]
```

```
A*B= [-1.   2.  -3.]
```

```
A/B= [-1.   2.  -3.]
```

# 张量运算



```
A = torch.randn((2,3))  
B = torch.randn((3,4))  
C = A@B  
print(A.shape)  
print(B.shape)  
print(C.shape)
```

```
torch.Size([2, 3])  
torch.Size([3, 4])  
torch.Size([2, 4])
```

# 张量运算



```
A = torch.tensor([[1,-1],[2,3],[4,5]])
```

```
B = A**2
```

```
print('A=',A.numpy())
```

```
print('B=',B.numpy())
```

```
A= [[ 1 -1]
```

```
 [ 2  3]
```

```
 [ 4  5]]
```

```
B= [[ 1  1]
```

```
 [ 4  9]
```

```
[16 25]]
```

# 张量运算



```
W = torch.randn((10,784))  
b = torch.randn((10,1))  
  
x = torch.randn((28,28)).reshape((784,1))  
y = W@x + b  
  
print(y)
```

```
tensor([[ 29.2908],  
        [-43.5128],  
        [-27.3771],  
        [ 27.4112],  
        [ 27.8641],  
        [ 26.4428],  
        [-36.5835],  
        [  5.1869],  
        [ 51.5410],  
        [-20.5552]])
```

# BroadCasting



```
x = torch.tensor([1,2,3])  
y = torch.tensor([[ -1],[ 0],[ 1]])  
u = x * y
```

```
u: tensor([[ -1, -2, -3],  
          [  0,  0,  0],  
          [  1,  2,  3]])
```

```
x: tensor([1, 2, 3])  
y: tensor([[ -1],  
          [  0],  
          [  1]])
```

```
x: torch.Size([3])  
y: torch.Size([3, 1])  
u: torch.Size([3, 3])
```

# BroadCasting



```
x = torch.tensor([1,2,3])
z = torch.ones((2,3,3),dtype = torch.int64)
v = x * z
```

```
x: tensor([1, 2, 3])
z: tensor([[[1, 1, 1],
            [1, 1, 1],
            [1, 1, 1]],

          [[1, 1, 1],
            [1, 1, 1],
            [1, 1, 1]]])
```

```
x: torch.Size([3])
z: torch.Size([2, 3, 3])
v: torch.Size([2, 3, 3])
```

```
v: tensor([[[[1, 2, 3],
             [1, 2, 3],
             [1, 2, 3]],

            [[1, 2, 3],
             [1, 2, 3],
             [1, 2, 3]]]])
```



# BroadCasting



```
y = torch.tensor([[[-1], [0], [1]]])
z = torch.ones((2,3,3), dtype = torch.int64)
w = y * z
```

```
y: tensor([[-1],
           [ 0],
           [ 1]])
z: tensor([[[1, 1, 1],
            [1, 1, 1],
            [1, 1, 1]],

           [[1, 1, 1],
            [1, 1, 1],
            [1, 1, 1]]])
```

```
y: torch.Size([3, 1])
z: torch.Size([2, 3, 3])
w: torch.Size([2, 3, 3])
```

```
w: tensor([[[[-1, -1, -1],
              [ 0,  0,  0],
              [ 1,  1,  1]],

            [[-1, -1, -1],
              [ 0,  0,  0],
              [ 1,  1,  1]]]])
```

# 例：彩色图像灰度化



```
from PIL import Image
img = np.asarray(Image.open('balloons.jpg'))
img_tensor = torch.tensor(img).to(dtype = torch.float32)
print(img_tensor.size())
```

```
torch.Size([605, 910, 3])
```

```
weights = torch.tensor([0.2126,0.7152,0.0722])
```

```
gray = torch.sum(weights * img_tensor,dim = 2)
```

```
torch.Size([605, 910])
```

```
img_tensor1 = img_tensor.permute([2,0,1])
```

```
torch.Size([3, 605, 910])
```

```
img_r,img_g,img_b = img_tensor1[0],img_tensor1[1],img_tensor1[2]
```

## 例2：图像标准化



```
img_files = ['balloons.jpg', 'happydog.jfif', 'sunflower.jpg', 'Woolsthorpe-Manor.jpg']
```

```
img_batch = torch.zeros(len(img_files), 3, 480, 640, dtype = torch.uint8)
```

```
for i, f in enumerate(img_files):
```

```
    img = np.array(Image.open(f).resize((640, 480)))
```

```
    img_tensor = torch.from_numpy(img)
```

```
    img_tensor = img_tensor.permute(2, 0, 1)
```

```
    img_batch[i] = img_tensor
```

```
img_batch : torch.Size([4, 3, 480, 640])
```

## 例2：图像标准化



```
img_batch = img_batch.float() / 255.0  
mean = torch.tensor([0.485, 0.456, 0.406])  
std = torch.tensor([0.229, 0.224, 0.225])
```

```
print('batch:',img_batch.size())  
print('mean:',mean.size())  
print('std:',std.size())
```

```
batch: torch.Size([4, 3, 480, 640])  
mean: torch.Size([3])  
std: torch.Size([3])
```

```
mean_unsqueezed = mean.unsqueeze(dim = 1).unsqueeze(dim=1)  
std_unsqueezed = std.unsqueeze(dim = 1).unsqueeze(dim = 1)  
print(mean_unsqueezed.size())
```

```
torch.Size([3, 1, 1])
```

```
img_batch_normalized = (img_batch - mean_unsqueezed) / std_unsqueezed
```

# 例3：使用张量编写神经元模型



```
class Neuron():  
    def __init__(self, in_features):  
        self.dim = in_features  
        self.W = torch.zeros((1,self.dim))  
        self.b = torch.zeros(1)  
  
    def __sigmoid__(self, z):  
        return 1/(1 + torch.exp(-z))  
  
    def __transfer__(self, x):  
        return self.W@x + self.b
```

# 例3：使用张量编写神经元模型



```
def __update__(self, dW,db,lr):  
    self.W = self.W + lr * dW  
    self.b = self.b + lr * db  
  
def __calc_loss__(self, Y,rho):  
    loss = -torch.log(rho[Y==1]).sum() - torch.log(1 - rho[Y==0]).sum()  
    loss = loss / Y.shape[0]  
    return loss
```

# 例3：使用张量编写神经元模型



```
def __backward__(self, Y, rho):  
    err = Y - rho  
    dW = err @ X.T/Y.shape[0]  
    db = err.mean()  
    return dW, db
```

```
def predict(self, x):  
    z = self.__transfer__(x)  
    rho = self.__sigmoid__(z)  
    return rho
```

# 例3：使用张量编写神经元模型



```
def fit(self, X,Y,max_iter=100,lr = 0.1):  
    n = X.shape[1]  
    assert(X.shape[0]==self.dim)  
    assert(n==Y.shape[0])  
  
    for iter in range(max_iter):  
        rho = self.predict(X).squeeze()  
        loss = self.__calc_loss__(Y,rho)  
        print('iter=',iter,',loss=',loss.item())  
        dW,db = self.__backward__(Y,rho)  
        self.__update__(dW,db,lr)
```



# 例3：使用张量编写神经元模型




```
def GenerateSamples(n):  
    x1 = torch.randn((2,n)) + 2  
    x2 = torch.randn((2,n)) - 2  
    y1 = torch.ones((n))  
    y2 = torch.zeros((n))  
    x = torch.cat((x1,x2),dim = 1)  
    y = torch.cat((y1,y2),dim = 0)  
  
    return x,y
```

# 例3：使用张量编写神经元模型



```
X,Y = GenerateSamples(30)
print(X.shape)
plt.plot(X[0,Y==1],X[1,Y==1], 'r+')
plt.plot(X[0,Y==0],X[1,Y==0], 'bo')

cell = Neuron(2)
cell.fit(X,Y,100,1)
```



```
iter= 0 , loss= 0.6931470632553101
iter= 1 , loss= 0.04438574239611626
iter= 2 , loss= 0.041226837784051895
iter= 3 , loss= 0.038746096193790436
iter= 4 , loss= 0.036735840141773224
iter= 5 , loss= 0.03506697714328766
iter= 6 , loss= 0.0336545892059803
iter= 7 , loss= 0.032440345734357834
```


# 例3：使用张量编写神经元模型



```
w1,w2 = cell.W[0,0],cell.W[0,1]
b = cell.b
minx = torch.min(X,dim = 1)
maxx = torch.max(X,dim = 1)

xs = torch.cat((minx[0],maxx[0]))
ys = -(xs*w1+b)/w2

plt.plot(X[0,Y==1],X[1,Y==1], 'r+')
plt.plot(X[0,Y==0],X[1,Y==0], 'b+')
plt.plot(xs,ys, 'k-')
```



# 例3：使用张量编写神经元模型



#测试集预测效果

```
X_test,Y_test = GenerateSamples(30)
```

```
Y_hat = torch.where(cell.predict(X_test).squeeze()>0.5,1,0)
```

```
plt.plot(X_test[0,Y_test==1],X_test[1,Y_test==1], 'r+')
```

```
plt.plot(X_test[0,Y_test==0],X_test[1,Y_test==0], 'b+')
```

```
plt.plot(X_test[0,Y_hat==1],X_test[1,Y_hat==1], 'ro',fillstyle='none')
```


```
plt.plot(X_test[0,Y_hat==0],X_test[1,Y_hat==0], 'bo',fillstyle='none')
```

```
acc = torch.mean((Y_hat==Y_test).to(torch.float32)).item()
```

```
print('Test Acc = ',acc)
```

```
Test Acc = 0.9833333492279053
```

# 例3：使用张量编写神经元模型



# 练习一：使用Tensor实现感知器模型

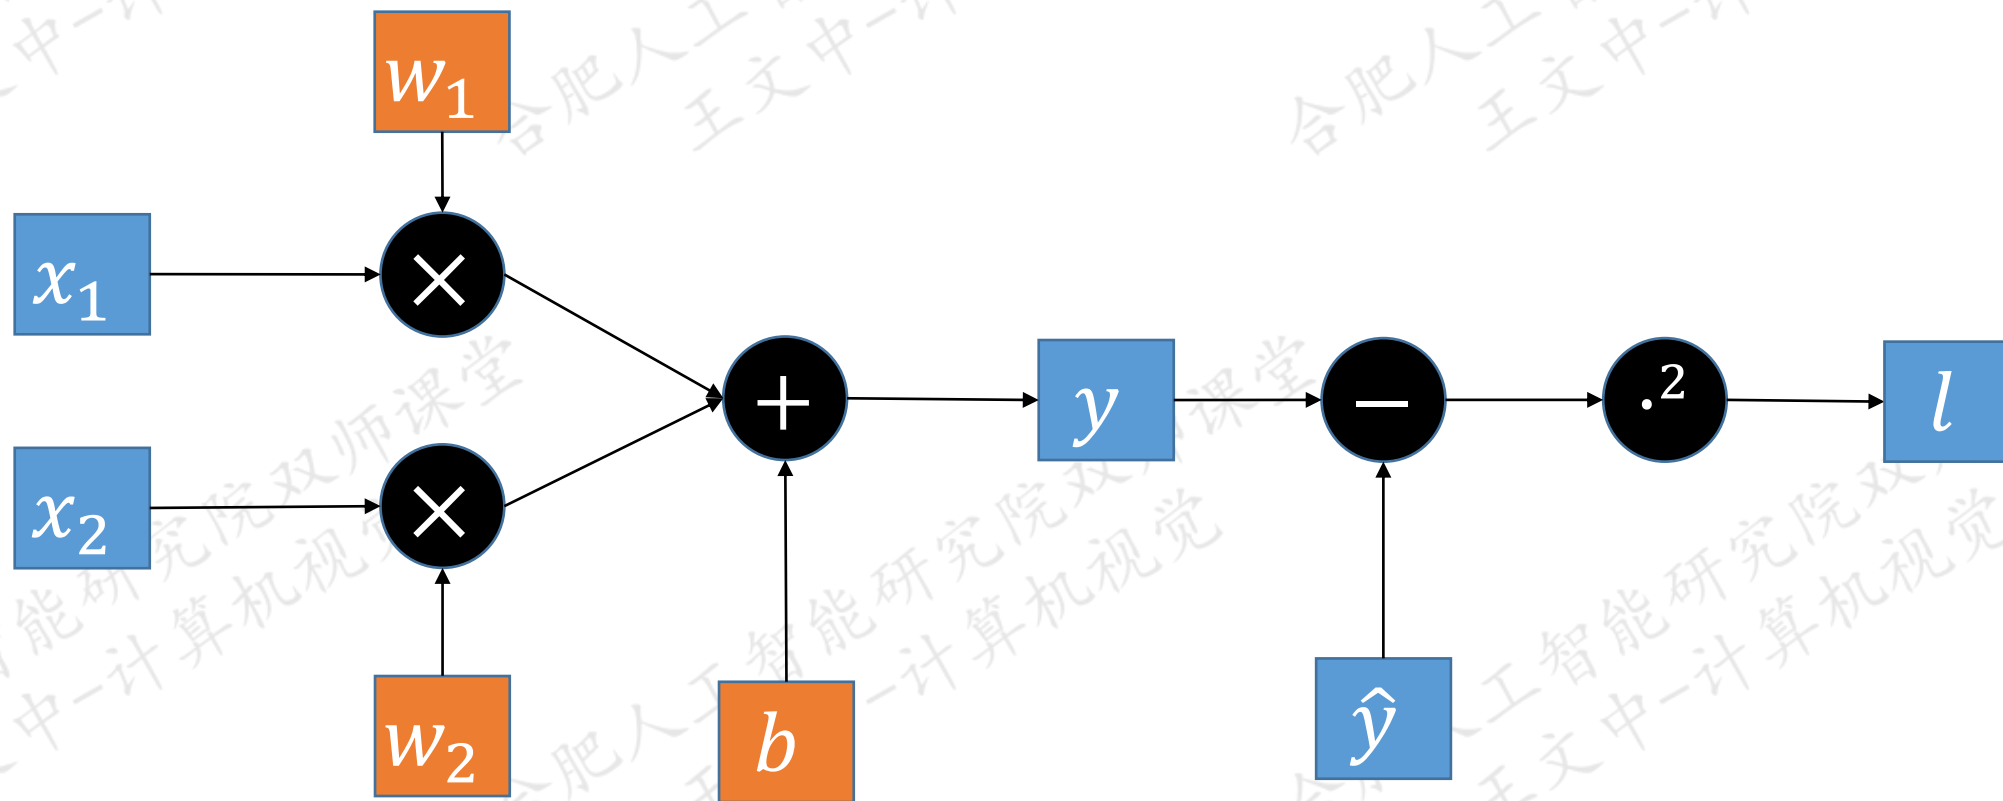


## 1.2 自动求导(AutoGrad)

# 计算图



$$y = h(x; w, b) = w_1 \times x_1 + w_2 \times x_2 + b$$
$$l = (y - \hat{y})^2$$





# 计算图



$$y = h(x; w, b) = w_1 \times x_1 + w_2 \times x_2 + b$$
$$l = (y - \hat{y})^2$$

$$t_1 = w_1 \times x_1$$

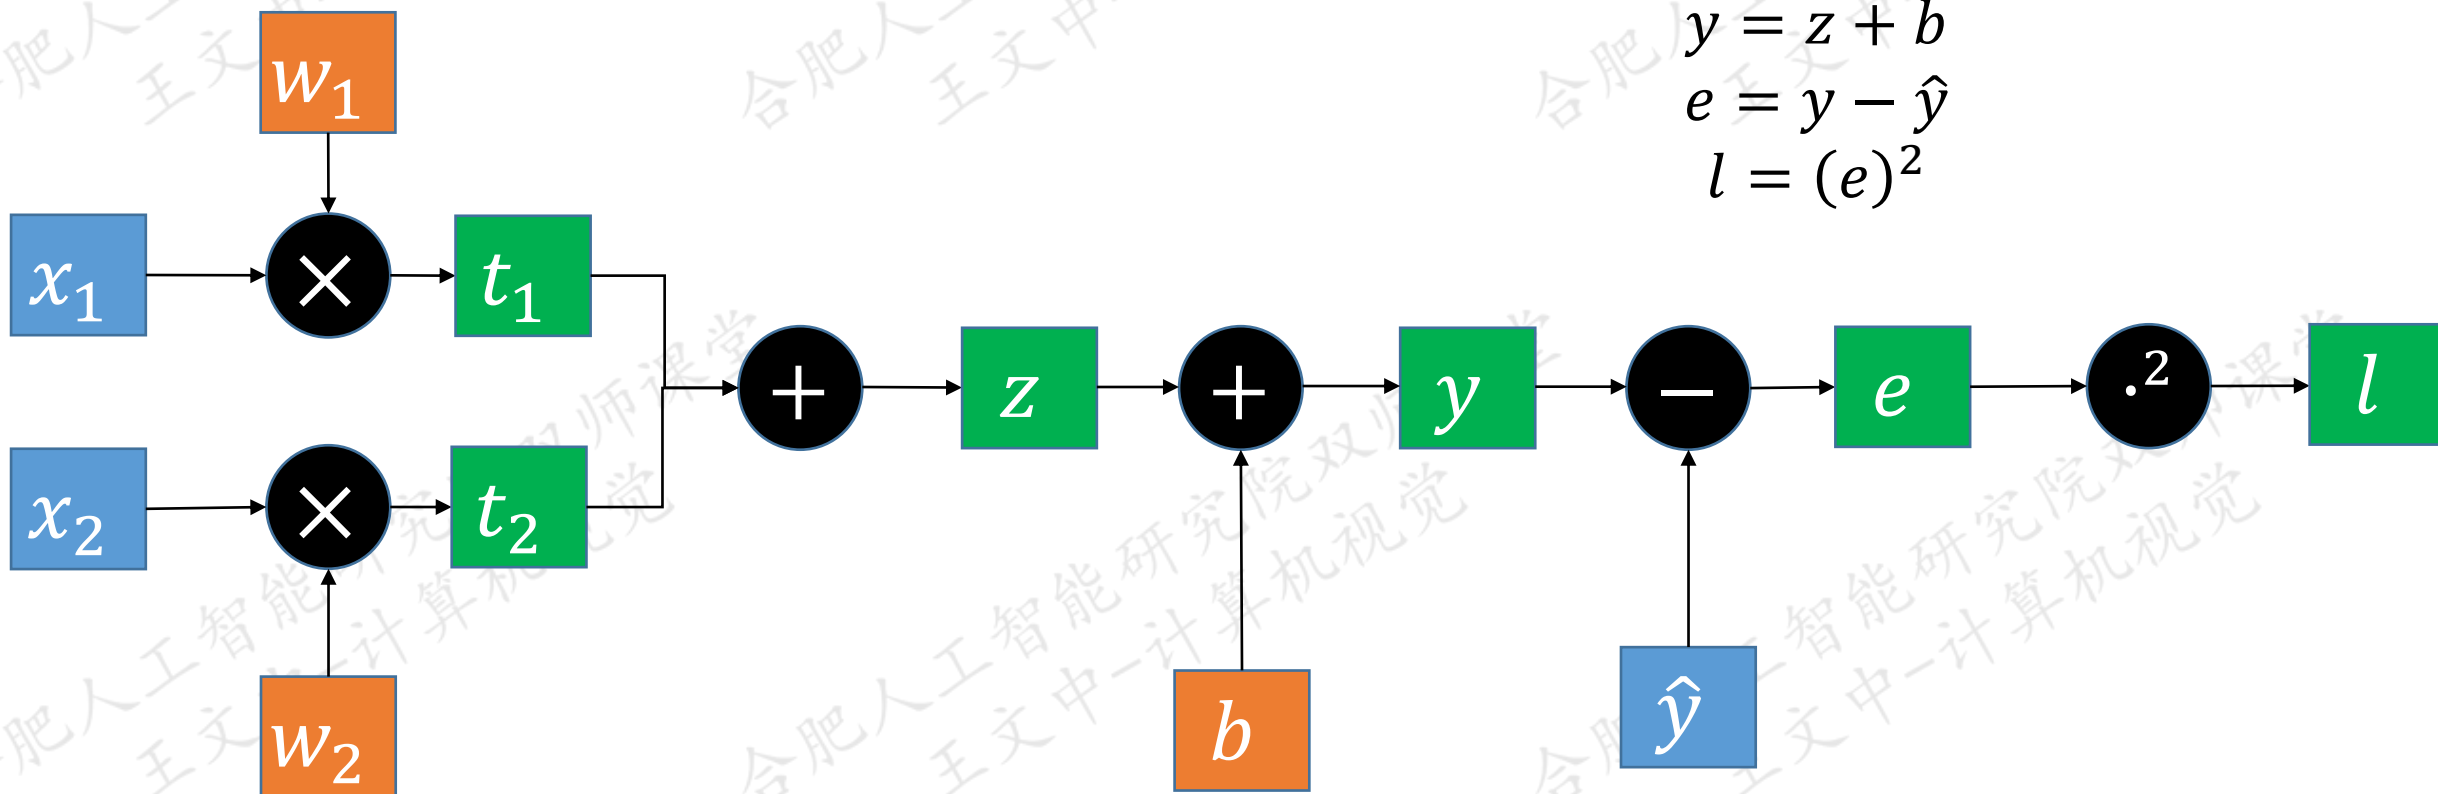
$$t_2 = w_2 \times x_2$$

$$z = t_1 + t_2$$

$$y = z + b$$

$$e = y - \hat{y}$$

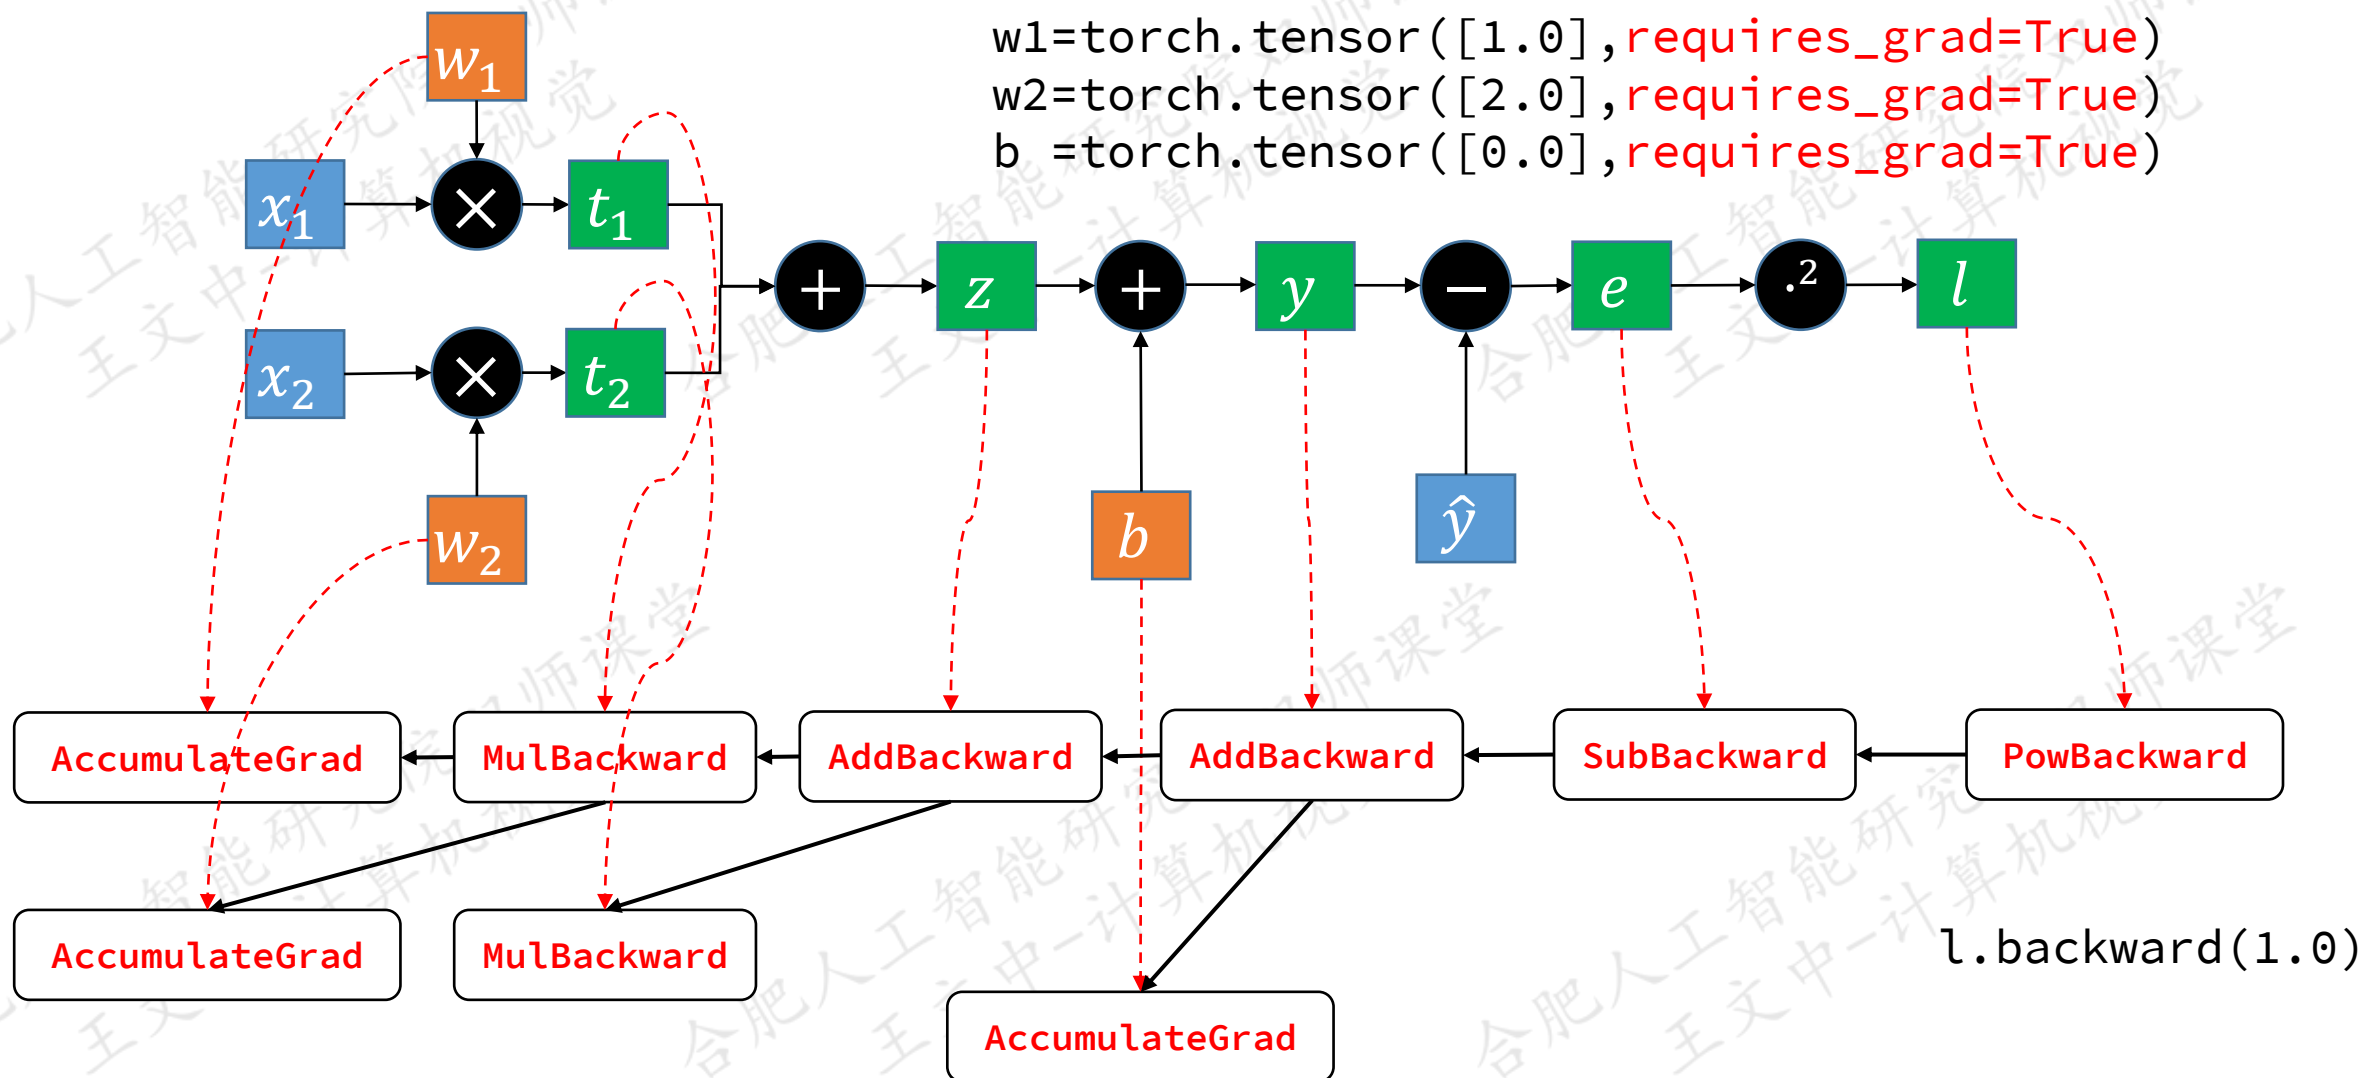
$$l = (e)^2$$



# 计算图



```
w1=torch.tensor([1.0],requires_grad=True)
w2=torch.tensor([2.0],requires_grad=True)
b =torch.tensor([0.0],requires_grad=True)
```



`l.backward(1.0)`

$$y = h(x; w, b) = w * x + b$$

$$loss = (y - \hat{y})^2$$

```
x = torch.tensor([1.0])
w = torch.tensor([1.0], requires_grad = True)
y = w*x
print(y)
```

```
tensor([1.], grad_fn=<MulBackward0>)
```

```
loss = (y - 2)**2
print(loss)
```

```
tensor([1.], grad_fn=<PowBackward0>)
```

```
loss.backward()
print(w.grad)
```

```
tensor([-2.])
```

```
with torch.no_grad():
    w -= 0.1*w.grad
print(w)
```

```
tensor([1.2000], requires_grad=True)
```

```
x = torch.tensor([1.0])
w = torch.tensor([1.0], requires_grad = True)
y = w*x

for epoch in range(3):
    print('epoch:%d'%(epoch))
    loss = (y - 2)**2
    loss.backward()
    print(w.grad)
```

```
epoch:0
tensor([-2.])
epoch:1
```

**RuntimeError:** Trying to backward through the graph a second time, but the buffers have already been freed. Specify retain\_graph=True when calling backward the first time.

```
x = torch.tensor([1.0])
w = torch.tensor([1.0], requires_grad = True)
y = w*x

for epoch in range(3):
    print('epoch:%d'%(epoch))
    loss = (y - 2)**2
    loss.backward(retain_graph=True)
    print(w.grad)
```


```
epoch:0
tensor([-2.])
epoch:1
tensor([-4.])
epoch:2
tensor([-6.])
```

```
x = torch.tensor([1.0])
w = torch.tensor([1.0], requires_grad = True)
y = w*x

for epoch in range(3):
    print('epoch:%d'%(epoch))
    if w.grad is not None:
        w.grad.zero_()
    loss = (y - 2)**2
    print(w.grad)
    loss.backward(retain_graph=True)
    print(w.grad)
```

```
epoch:0
None
tensor([-2.])
epoch:1
tensor([0.])
tensor([-2.])
epoch:2
tensor([0.])
tensor([-2.])
```

# 例1：一个简单的一元线性回归问题



#生成训练数据

`x = torch.rand(10)*10 - 5` #[-5,5]上的均匀分布随机数

`y = 2*x + 1 + torch.randn(10)` #受随机噪声污染的y

#定义回归模型

```
def linearReg(x,w,b):
```

```
    y = w * x + b
```

```
    return y
```

#定义损失函数

```
def lossFn(y,y_hat):
```

```
    squared_errors = (y - y_hat)**2
```

```
    loss = squared_errors.mean()
```

```
    return loss
```



#定义训练函数

```
def trainModel(x,y,w,b,epochs,lr = 0.1):  
    for epoch in range(1, epochs + 1):  
        if w.grad is not None:  
            w.grad.zero_()  
        if b.grad is not None:  
            b.grad.zero_()  
  
        y_pred = linearReg(x,w,b)  
  
        loss = lossFn(y_pred,y)  
  
        loss.backward()  
  
        with torch.no_grad():  
            w -= lr * w.grad  
            b -= lr * b.grad  
  
        print('Epoch = %d, Loss = %f, w = %f, b = %f'%(  
            epoch,loss.detach().numpy(),w.detach().numpy(),b.detach().numpy()))  
    return w,b
```

#训练模型

#初始化参数w,b

```
w = torch.tensor([0.0],requires_grad = True)
```

```
b = torch.tensor([0.0],requires_grad = True)
```

```
w,b = trainModel(x,y,w,b,epochs = 10, lr = 0.2)
```

```
Epoch = 1, Loss = 12.116911, w = 2.548692,b = 0.003886
```

```
Epoch = 2, Loss = 4.572442, w = 1.183241,b = 0.532312
```

```
Epoch = 3, Loss = 2.037711, w = 2.024282,b = 0.567515
```

```
Epoch = 4, Loss = 1.181779, w = 1.580700,b = 0.762241
```

```
Epoch = 5, Loss = 0.891036, w = 1.858682,b = 0.787515
```

```
Epoch = 6, Loss = 0.791604, w = 1.714883,b = 0.860059
```


```
Epoch = 7, Loss = 0.757336, w = 1.806942,b = 0.873903
```

```
Epoch = 8, Loss = 0.745425, w = 1.760451,b = 0.901212
```

```
Epoch = 9, Loss = 0.741245, w = 1.791010,b = 0.908000
```

```
Epoch = 10, Loss = 0.739764, w = 1.776030,b = 0.918381
```

```
y_pred = linearReg(x,w,b)
plt.plot(x.numpy(),y.numpy(),'*')
plt.plot(x.numpy(),2*x.numpy()+1,'r-')
plt.plot(x.numpy(),y_pred.detach().numpy(),'ko')
plt.show()
```



# 使用torch.optim训练模型







#定义训练函数

```
def trainModel(x,y,w,b,epochs,lr = 0.1):  
    for epoch in range(1, epochs + 1):  
        if w.grad is not None:  
            w.grad.zero_()  
        if b.grad is not None:  
            b.grad.zero_()  
  
        y_pred = linearReg(x,w,b)  
  
        loss = lossFn(y_pred,y)  
  
        loss.backward()  
  
        with torch.no_grad():  
            w -= lr * w.grad  
            b -= lr * b.grad  
  
    return w,b
```

import torch.optim as optim

```
def trainModel(x,y,w,b,epochs,lr = 0.1):  
    optimizer = optim.SGD([w,b],lr = lr)  
    for epoch in range(1, epochs + 1):  
        optimizer.zero_grad()  
        y_pred = linearReg(x,w,b)  
        loss = lossFn(y_pred,y)  
        loss.backward()  
        optimizer.step()  
  
    return w,b
```

# 使用torch.optim训练模型



# 练习二



- 使用Tensor和自动求导编写线性回归器

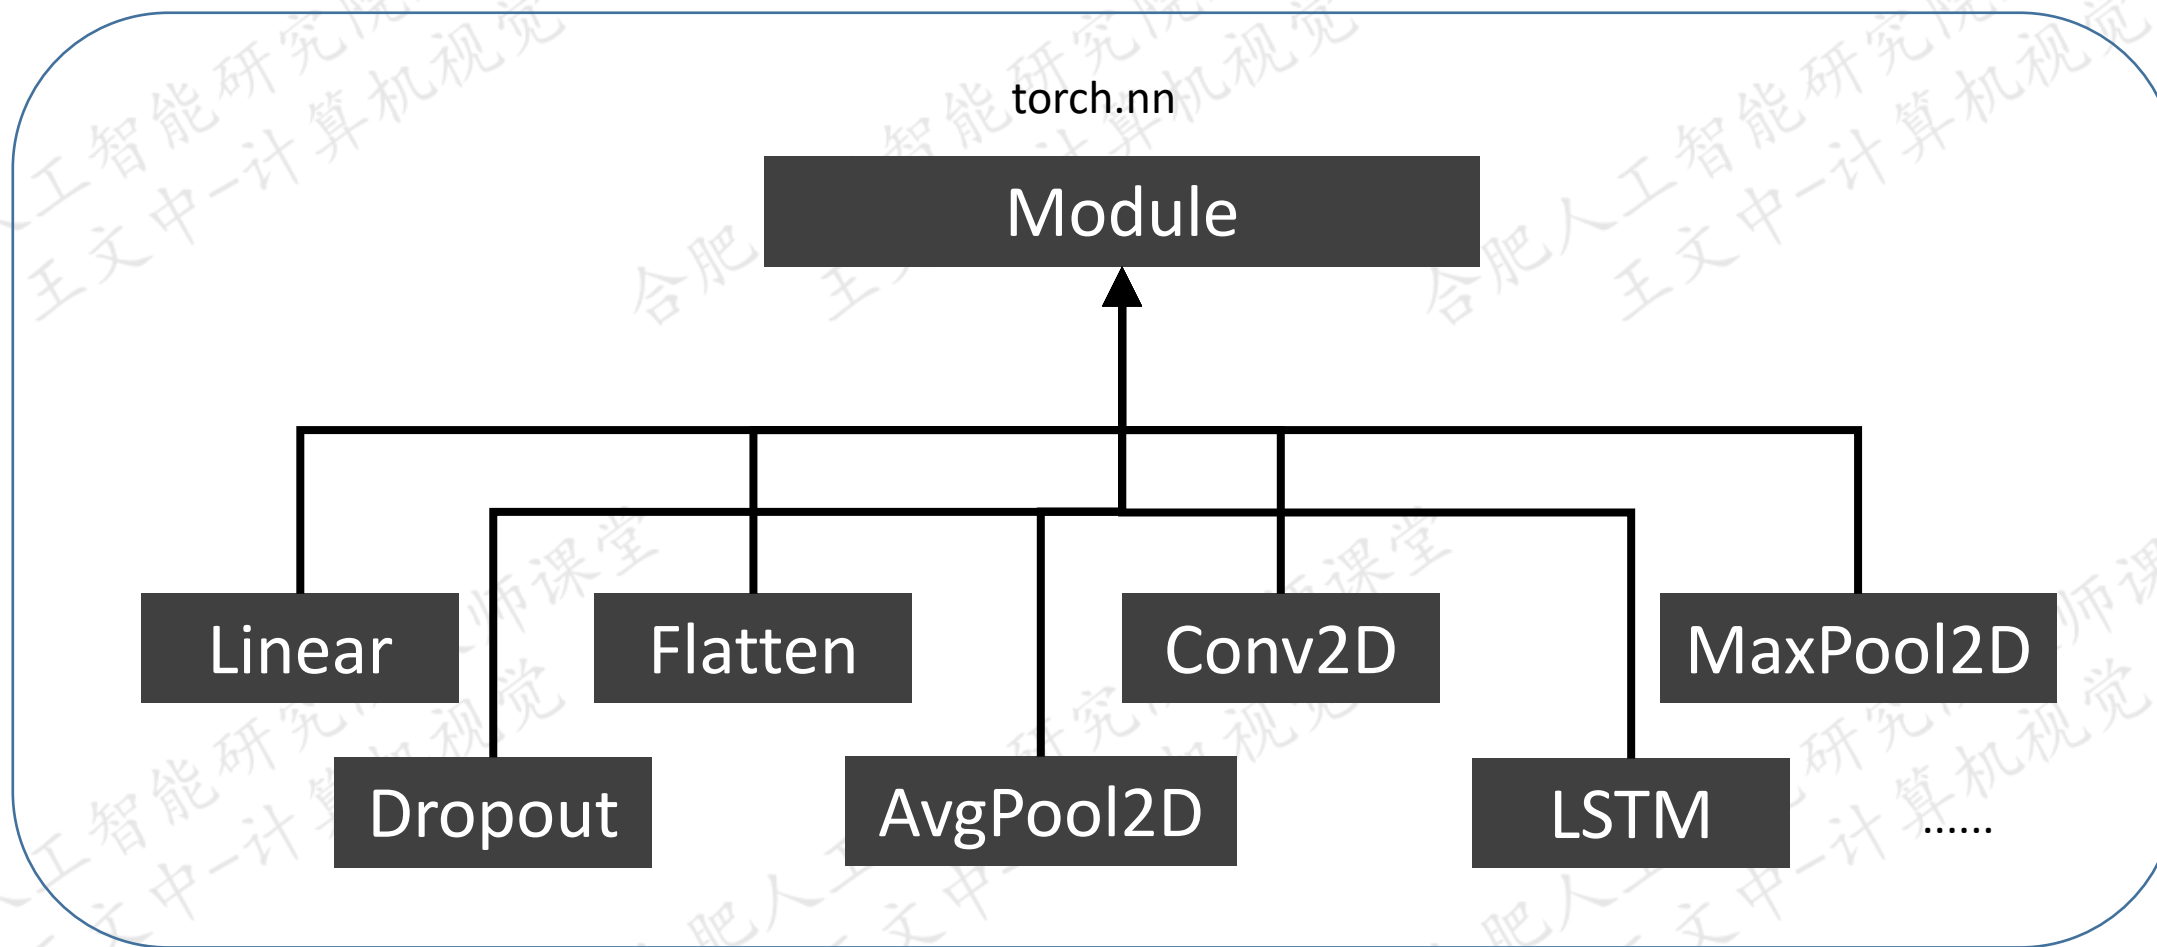
# 2. 多层感知器编程

torch.nn

- 编写网络模型
  - `torch.nn.Sequential`: 编写简单的模型
  - `torch.nn.Module`: 编写复杂的模型
- 训练模型:
  - 数据加载器
  - 损失函数: `torch.nn`
  - 优化器: `torch.nn.optim`



# torch.nn.Module



# 2.1 实现简单的模型

# 线性回归




```
torch.nn.Linear(in_features: int, out_features: int, bias: bool = True)
```

```
import torch.nn as nn
linReg = nn.Linear(1,1)

x = torch.rand(10)*10 - 5
y = 2*x + 1 + torch.randn(10)

y_pred = linReg(x.view(10,1))

plt.plot(x.numpy(),y.numpy(),'*')
plt.plot(x.numpy(),2*x.numpy()+1,'r-')
plt.plot(x.numpy(),y_pred.detach().numpy(),'ko')
plt.show()
```



# 线性回归



```
torch.nn.MSELoss(size_average=None, reduce=None, reduction: str = 'mean')
```

```
linReg = nn.Linear(1,1)
loss_fn = nn.MSELoss(reduction='mean')
linReg.train()
optimizer = optim.SGD(params = linReg.parameters(),lr = 0.1,momentum = 0.9)
for epoch in range(10):
    optimizer.zero_grad()
    y_pred = linReg(x.view(10,1))
    loss = loss_fn(y_pred, y)
    loss.backward()
    optimizer.step()
```


```
x_train,y_train = readData('circle-train.csv')
x = torch.tensor(x_train,dtype = torch.float32)
y = torch.tensor(y_train[:,np.newaxis],dtype = torch.float)

logReg = nn.Linear(2,1)
loss_fn = nn.BCEWithLogitsLoss()

logReg.train()
optimizer = optim.SGD(params = logReg.parameters(),lr = 0.1,momentum = 0.9)
for epoch in range(10):
    optimizer.zero_grad()
    y_pred = logReg(x)
    loss = loss_fn(y_pred, y)
    loss.backward()
    optimizer.step()
```

# 2.2 使用Sequential构造神经网络模型

# 构造序列式神经网络模型



# torch.nn.Sequential





```
model = nn.Sequential(  
    nn.Conv2d(1, 20, 5),  
    nn.ReLU(),  
    nn.Conv2d(20, 64, 5),  
    nn.ReLU()  
)
```



```
from collections import OrderedDict  
  
model = nn.Sequential(OrderedDict([  
    ('conv1', nn.Conv2d(1, 20, 5)),  
    ('relu1', nn.ReLU()),  
    ('conv2', nn.Conv2d(20, 64, 5)),  
    ('relu2', nn.ReLU())  
]))
```



# MNIST手写体识别



# MNIST手写体识别:设计一个单隐层神经网络



$$h = \sigma(W_h x + b_h)$$

$$p = \text{Softmax}(W_o h + b_o)$$

# torch.nn.Linear



```
torch.nn.Linear(in_features, out_features, bias=True,  
device=None, dtype=None)
```

$$out = W \times input + b$$

# torch.nn.Linear



```
x = torch.tensor([1.0,2.0,3.0])
lin = torch.nn.Linear(3,2)
y = lin(x)
print(y)
```

```
tensor([-0.7312,  0.5305], grad_fn=<AddBackward0>)
```

```
for v in lin.parameters():
    print(v)
```

```
Parameter containing:
tensor([[[-0.0212,  0.2302, -0.4544],
         [ 0.5579,  0.4064, -0.2442]], requires_grad=True)
Parameter containing:
tensor([ 0.1929, -0.1078], requires_grad=True)
```

# torch.nn.Softmax



```
torch.nn.Softmax(dim=None)
```

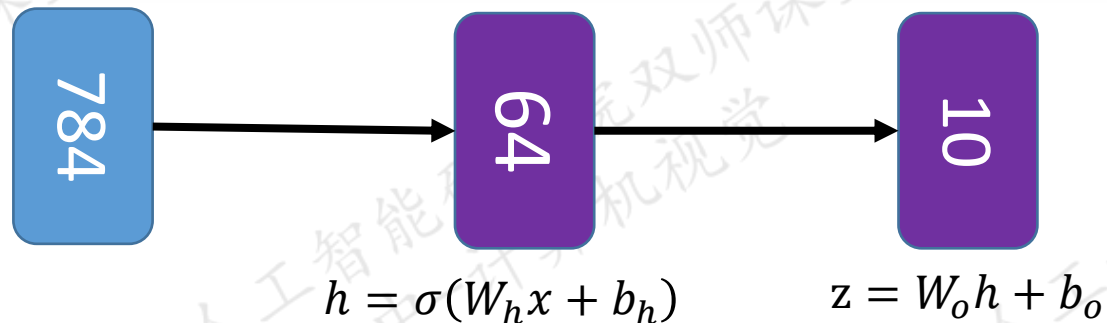
```
x = torch.tensor([[1.0, 3.0, 5.0], [4.0, 5.0, 1.0]])  
prob1 = torch.nn.Softmax(dim=0)  
y1 = prob1(x)
```

```
tensor([[0.0474, 0.1192, 0.9820],  
        [0.9526, 0.8808, 0.0180]])
```

```
prob2 = torch.nn.Softmax(dim=1)  
y2 = prob2(x)
```

```
tensor([[0.0159, 0.1173, 0.8668],  
        [0.2654, 0.7214, 0.0132]])
```

# MNIST手写体识别：构造神经网络



```
mnist_mlp = nn.Sequential(nn.Linear(784, 64),  
                           nn.Sigmoid(),  
                           nn.Linear(64, 10))  
  
optimizer = optim.SGD(mnist_mlp.parameters(), lr = 0.01)  
  
loss_fn = nn.CrossEntropyLoss()
```

# torch.nn.CrossEntropyLoss



$$l(\theta; z, y) = -\log\left(\frac{e^{z_y}}{\sum_{j=1}^k e^{z_j}}\right), z = h(x; \theta) = \begin{pmatrix} z_1 \\ z_2 \\ \vdots \\ z_k \end{pmatrix}$$

torch.nn.NLLLoss

$$l(\theta; z, y) = -z_y, y \in \{1, 2, \dots, k\}$$

$$z_j = \log(\rho_j), \rho_j = \frac{e^{z_j}}{\sum_{i=1}^k e^{z_i}}$$

# MNIST手写体识别：数据集与加载器



```
from torchvision import datasets, transforms

data_path = '../data/'
mnist = datasets.MNIST(data_path, download=True,
                       transform = transforms.ToTensor())

img, label = mnist[0]
img.shape                                torch.Size([1, 28, 28])
img = img.view(-1)
img.shape                                torch.Size([784])

mnist_loader = torch.utils.data.DataLoader(mnist, batch_size = 64,
                                           shuffle = True)
```




# MNIST手写体识别：数据集与加载器



```
mnist_loader = torch.utils.data.DataLoader(mnist,
                                             batch_size = 64,
                                             shuffle = True)

for i,data in enumerate(mnist_loader):
    print(data[0].shape)
    print(data[1].shape)
    break
```

```
torch.Size([64, 1, 28, 28])
torch.Size([64])
```



# MNIST手写体识别：训练神经网络



```
epochs = 10
mnist_mlp.train()
for epoch in range(epochs):
    for imgs, labels in mnist_loader:
        batch_size = imgs.shape[0]
        logits = mnist_mlp(imgs.view(batch_size, -1))
        loss = loss_fn(logits, labels)
        optimizer.zero_grad()
        loss.backward()
        optimizer.step()
```

# MNIST手写体识别：测试模型



```
test_loader = torch.utils.data.DataLoader(mnist_test,
                                          batch_size = 100, shuffle = False)

mnist_mlp.eval()
correct = 0
total = 0
with torch.no_grad():
    for imgs, labels in test_loader:
        batch_size = imgs.shape[0]
        logits = mnist_mlp(imgs.view(batch_size, -1))
        _, predicted = torch.max(logits.data, 1)
        total += batch_size
        correct += (predicted == labels).sum().item()
acc = correct / total
```

# 2.2 子类化 torch.nn.Module

# 子类化torch.nn.Module



- 从torch.nn.Module派生一个子类
- 编写\_\_init\_\_函数
  - 初始化对象
  - 要调用父类的构造函数\_\_init\_\_
- 编写forward函数
  - 构造网络

# 使用nn.Module构造网络模型



```
class mlp(nn.Module):  
  
    def __init__(self, in_dim, out_dim):  
        super(mlp, self).__init__()  
        self.lin1 = nn.Linear(in_features=in_dim, out_features=64)  
        self.relu = nn.ReLU(inplace=True)  
        self.lin2 = nn.Linear(in_features=64, out_features=out_dim)  
  
    def forward(self, x):  
        x = self.lin1(x)  
        x = self.relu(x)  
        x = self.lin2(x)  
  
        return x
```

```
mnist_mlp = mlp(784,10)
print(mnist_mlp)
```

```
mlp(
    (lin1): Linear(in_features=784, out_features=64, bias=True)
    (relu): ReLU(inplace)
    (lin2): Linear(in_features=64, out_features=10, bias=True)
)
```

```
x = torch.randn((1,784),dtype = torch.float32)
y = F.softmax(mnist_mlp(x),dim=1)
print(y.detach().numpy())
```

```
[[0.07569709 0.11641571 0.12871902 0.10927621 0.05882455 0.12559932
 0.11405984 0.07417176 0.0880156 0.1092209 ]]
```

```
print(mnist_mlp._modules)
```


```
OrderedDict([('lin1', Linear(in_features=784, out_features=64,  
bias=True)), ('relu', ReLU(inplace)), ('lin2', Linear(in_features=64,  
out_features=10, bias=True))])
```

```
x = torch.randn((1,784),dtype = torch.float32)  
y = mnist_mlp._modules['lin1'](x)  
print(y.detach().numpy())
```

```
[[-0.01645807  0.12257674 -0.3842733 -0.10875969  0.16210961 -0.9262785  
  0.3560028 -0.3479771 -0.4386911 -0.37093288 -0.0556463 -0.4703986  
  0.54101855 -0.51167864  0.08044843  0.21852194  1.1178634  0.22886397  
 -0.6046644 -0.36042506 -0.0121364 -0.37140787  0.41181365  0.9680239  
 -0.44987577  0.11921623 -0.55106914  0.18517387  0.56142396 -0.59405994  
  0.47401157  1.1164094  0.48247746  0.11364826  0.91904306  0.9782668  
 -0.81544924  0.2937653 -1.3910329 -0.10774229 -0.03324178  0.54723066  
 -0.43887684  0.12371415  0.03490007  0.62052786 -0.10143616 -0.43207315  
  0.42250878  0.49805558 -0.69287837 -0.2873435 -0.5114934  0.6561675  
  1.0234826  0.04920949  0.3638496  0.5078668  0.1270785 -0.18455932  
 -1.1579272 -0.6274739 -0.23763435 -0.7309237 ]]
```



# 用Module实现复杂的模型



# 用Module实现复杂的模型




```
import torch.nn.functional as F

class mlp(torch.nn.Module):
    def __init__(self, in_dim, out_dim):
        super(mlp, self).__init__()
        self.lin1 = nn.Linear(in_features=in_dim, out_features=128)
        self.lin2 = nn.Linear(in_features=128, out_features=128)
        self.lin3 = nn.Linear(in_features=128, out_features=out_dim)

    def forward(self, x):
        x1 = F.relu(self.lin1(x))
        x2 = F.relu(self.lin2(x1))
        x3 = x1+x2
        x4 = self.lin3(x3)
        return x4
```

# 用Module实现复杂的模型



# 用Module实现复杂的模型




```
import torch.nn.functional as F

class mlp(nn.Module):
    def __init__(self, in_dim, out_dim):
        super(mlp, self).__init__()
        self.lin1 = nn.Linear(in_features=in_dim, out_features=128)
        self.lin2 = nn.Linear(in_features=128, out_features=128)
        self.lin3 = nn.Linear(in_features=256, out_features=out_dim)

    def forward(self, x):
        x1 = F.relu(self.lin1(x))
        x2 = F.relu(self.lin2(x1))
        x3 = torch.cat((x1, x2), dim = 1)
        x4 = self.lin3(x3)
        return x4
```

# 用Module实现复杂的模型



$$h_1 = \text{relu}(W_{h1}x + b_{h1}) \quad h_2 = \text{relu}(W_{h2}h_1 + b_{h2}) \quad z = \text{Softmax}(W_o h_2 + b_o) \\ \|h_2\| = 1$$

$$\text{contrastive loss}(x_1, x_2, y) = y \times d(x_1, x_2) + (1 - y) \times \max(\alpha - d(x_1, x_2), 0)$$

# 用Module实现复杂的模型



```
def contrastive_loss(x,y, margin = 0.3):  
    pair_dist = torch.cdist(x,features)  
  
    l1 = torch.unsqueeze(y,dim=0)  
    l2 = torch.unsqueeze(y,dim=1)  
    pair_mask = torch.logical_not(torch.eye(y.shape[0],dtype=torch.bool))  
  
    pos_mask = torch.logical_and(l1==l2, pair_mask)  
    neg_mask = torch.logical_and(l1!=l2, pair_mask)  
  
    pos_dist = torch.masked_select(pair_dist,pos_mask)  
    neg_dist = torch.masked_select(pair_dist,neg_mask)  
  
    pos_loss = torch.mean(pos_dist)  
    margin = torch.tensor(margin)  
    neg_loss = torch.mean(torch.max(margin - neg_dist,torch.tensor(0.0)))  
    return pos_loss + neg_loss
```

# 用Module实现复杂的模型



```
labels = torch.tensor([0,1,2,0,2])
l1 = torch.unsqueeze(labels,dim=0)
l2 = torch.unsqueeze(labels,dim=1)
```

```
l1:tensor([[0, 1, 2, 0, 2]])
```

```
l2:tensor([[0],
          [1],
          [2],
          [0],
          [2]])
```

```
pair_mask = torch.logical_not(torch.eye(labels.shape[0],dtype=torch.bool))
pos_mask = torch.logical_and(l1==l2,pair_mask)
neg_mask = torch.logical_and(l1!=l2,pair_mask)
```

```
sim = torch.randn((5,5))
```

```
tensor([[ 0.5344,  0.7295, -0.4729,  0.4417,  1.5117],
        [ 0.3096, -0.1728,  0.4240, -1.8273,  0.6577],
        [-0.0897, -0.8336, -0.2095,  0.4435,  1.1995],
        [-0.5547,  0.6684, -1.2100, -0.2603,  0.6245],
        [ 0.3822,  1.2045,  0.2586,  0.9755, -0.2574]])
```

```
pos_sim = torch.masked_select(sim,pos_mask)
neg_sim = torch.masked_select(sim,neg_mask)
```

```
tensor([ 0.4417,  1.1995, -0.5547,  0.2586])
tensor([ 0.7295, -0.4729,  1.5117,  0.3096,  0.4240, -1.8273,  0.6577, -0.0897,
        -0.8336,  0.4435,  0.6684, -1.2100,  0.6245,  0.3822,  1.2045,  0.9755])
```

```
tensor([[False,  True,  True,  True,  True],
        [ True, False,  True,  True,  True],
        [ True,  True, False,  True,  True],
        [ True,  True,  True, False,  True],
        [ True,  True,  True,  True, False]])
```

```
tensor([[False, False, False,  True, False],
        [False, False, False, False, False],
        [False, False, False, False,  True],
        [ True, False, False, False, False],
        [False, False,  True, False, False]])
```

```
tensor([[False,  True,  True, False,  True],
        [ True, False,  True,  True,  True],
        [ True,  True, False,  True, False],
        [False,  True,  True, False,  True],
        [ True,  True, False,  True, False]])
```

# 用Module实现复杂的模型



```
class mlp(nn.Module):
    def __init__(self, in_features):
        super(mlp, self).__init__()

        self.fnet = nn.Sequential(OrderedDict([
            ('lin1', nn.Linear(in_features, 256)),
            ('relu1', nn.ReLU()),
            ('lin2', nn.Linear(256, 128)),
            ('relu2', nn.ReLU())]))

        self.fc = nn.Linear(128, 10)
        self.softmax = nn.Softmax(dim=1)

    def forward(self, x):
        f = F.normalize(self.fnet(x), dim=1)
        logits = self.fc(f)
        p = self.softmax(logits)

        return f, logits, p
```



# 用Module实现复杂的模型



```
def train(model, data_loader, lr=0.01):
    epochs = 10
    model.train()
    optimizer = optim.RMSprop(model.parameters(), lr = lr)

    identity_loss = torch.nn.CrossEntropyLoss()

    for epoch in range(epochs):
        for imgs, labels in data_loader:
            batch_size = imgs.shape[0]
            features, logits, prob = model(imgs.view(batch_size, -1))

            contr_loss = contrastive_loss(features, labels)
            ident_loss = identity_loss(logits, labels)
            loss = contr_loss + ident_loss

            optimizer.zero_grad()
            loss.backward()
            optimizer.step()
```

## 使用pytorch开发神经网络模型的基本步骤:

- 1. 构造模型:
  - torch.nn.Sequential类
  - torch.nn.Module类
- 2. 训练模型:
  - 损失函数: torch.nn模块
  - 优化算法: torch.optim模块
  - 数据加载: torch.utils.data模块

# 练习三



## • 1. 糖尿病预测（回归）

```
from sklearn.datasets import load_diabetes
diabetes_dataset = load_diabetes()
data = diabetes_dataset['data']
targets = diabetes_dataset['target']
print(data.shape)
print(targets.shape)
```


```
(442, 10)
(442,)
```

## • 2. MNIST手写体识别（分类）

# 3.卷积神经网络编程

# 3.1 编写卷积网络模型

# 卷积神经网络



AlexNet

`torch.nn.Conv2d`

`torch.nn.MaxPool2d`

`torch.nn.AvgPool2d`

```
torch.nn.Conv2d(in_channels: int,  
out_channels: int,  
kernel_size: Union[int, Tuple[int, int]],  
stride: Union[int, Tuple[int, int]] = 1,  
padding: Union[int, Tuple[int, int]] = 0,  
dilation: Union[int, Tuple[int, int]] = 1,  
groups: int = 1,  
bias: bool = True,  
padding_mode: str = 'zeros')
```

$input: N \times C_{in} \times H_{in} \times W_{in}$   
 $output: N \times C_{out} \times H_{out} \times W_{out}$

```
m = nn.Conv2d(16, 33, 3, stride=2)  
m = nn.Conv2d(16, 33, (3, 5), stride=(2, 1), padding=(4, 2))  
m = nn.Conv2d(16, 33, (3, 5), stride=(2, 1), padding=(4, 2), dilation=(3, 1))
```

# nn.MaxPool2d



```
torch.nn.MaxPool2d(kernel_size,  
stride=None,  
padding=0,  
dilation=1,  
return_indices=False,  
ceil_mode=False)
```


```
m1 = nn.MaxPool2d(2, 2)  
m2 = nn.MaxPool2d((4, 2), stride=(2, 1), padding=(2, 1))  
x = torch.rand((1,12,12))#x.dim==3 or 4  
y1 = m1(x)      torch.Size([1, 6, 6])  
y2 = m2(x)      torch.Size([1, 7, 13])
```



# CIFAR10



airplane



automobile



bird



cat



deer



dog



frog



horse



ship

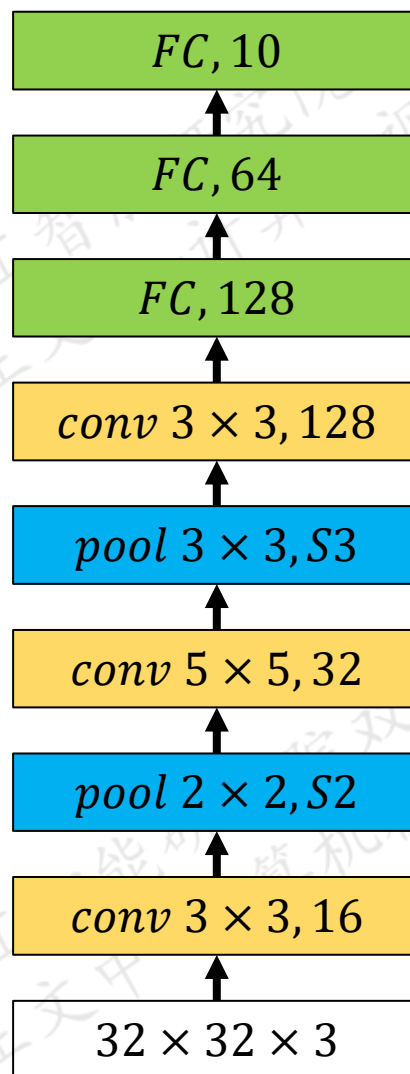


truck



$32 \times 32 \times 3$

# CIFAR10



```
class CNN(nn.Module):
```

```
    def __init__(self):
```

```
        #注意：首先调用父类的初始化函数
```

```
        super(CNN, self).__init__()
```

```
        #定义卷积、池化以及全连接操作
```

conv  $3 \times 3, 16$

pool  $2 \times 2, S2$

conv  $5 \times 5, 32$

pool  $3 \times 3, S3$

conv  $3 \times 3, 128$

FC, 128

FC, 64

FC, 10

```
        self.conv1 = nn.Conv2d(in_channels=3, out_channels=16, kernel_size=3)
```

```
        self.pool1 = nn.MaxPool2d(kernel_size=2, stride=2)
```

```
        self.conv2 = nn.Conv2d(in_channels=16, out_channels=32, kernel_size=5)
```

```
        self.pool2 = nn.MaxPool2d(kernel_size=3, stride=3)
```

```
        self.conv3 = nn.Conv2d(in_channels=32, out_channels=128, kernel_size=3)
```

```
        self.fc1 = nn.Linear(128, 128)
```

```
        self.fc2 = nn.Linear(128, 64)
```

```
        self.fc3 = nn.Linear(64, 10)
```

```
def forward(self, x):
```

```
    #在前向函数中构造出卷积网络
```

```
    #注意这里的x把不同层连接起来
```

```
    x = self.pool1(F.relu(self.conv1(x)))
```

```
    x = self.pool2(F.relu(self.conv2(x)))
```

```
    x = F.relu(self.conv3(x))
```

```
    #使用torch.Tensor.view函数，把一个多维张量拉直为一个1维张量（向量）
```

```
    x = x.view(-1, 128)
```

```
    #全连接层
```

```
    x = F.relu(self.fc1(x))
```

```
    x = F.relu(self.fc2(x))
```

```
    x = self.fc3(x)
```

```
    return x
```

# 组织训练样本datasets.ImageFolder



```
data/dog/0001.jpg  
data/dog/0002.jpg  
data/dog/harris.png  
.....  
data/cat/0001.png  
data/cat/ricky.png  
data/cat/adam_1.png
```

```
data_set = torchvision.datasets.ImageFolder(root='./data')
```

# 样本预处理transforms



```
normalize = transforms.Normalize((0.5, 0.5, 0.5), (0.5, 0.5, 0.5))
```

```
data_transform = transforms.Compose([  
    transforms.RandomResizedCrop(224),  
    transforms.RandomHorizontalFlip(),  
    transforms.ToTensor(), #把图像变换为张量  
    normalize, #注意规范化要在ToTensor之后  
)
```

```
data_set = torchvision.datasets.ImageFolder(root='./data',  
                                             transform=data_transform)
```

# 数据加载器torch.utils.data.DataLoader



```
train_loader = torch.utils.data.DataLoader(dataset=train_set,  
                                             batch_size=32,  
                                             shuffle=True,  
                                             num_workers=0)  
  
test_loader = torch.utils.data.DataLoader(dataset=test_set,  
                                           batch_size=32,  
                                           shuffle=False,  
                                           num_workers=0)
```



```
net = CNN()
if torch.cuda.is_available():
    device = torch.device("cuda:0")
else:
    device = torch.device("cpu")

net.to(device)
xentropy = nn.CrossEntropyLoss()
optimizer = optim.SGD(net.parameters(), lr=0.01, momentum=0.9)
net.train()
```



```

for epoch in range(50):
    running_loss = 0.0
    for i, data in enumerate(train_loader):
        inputs, labels = data[0].to(device), data[1].to(device)
        optimizer.zero_grad()
        outputs = net(inputs)
        loss = xentropy(outputs, labels)
        loss.backward()
        optimizer.step()
        running_loss += loss.item()

    if i % 10 == 9:
        print('[%d, %5d] loss = %.3f' % (epoch+1, i+1, running_loss/10))
        running_loss = 0.0

```

*input:  $N \times C \times H \times W$*   
*output:  $N \times K$*

## 3.2 使用预训练模型

# 保存模型



```
model = CNN()  
state_dict = model.state_dict()  
print('State_dict of CNN model:')  
for param_tensor in state_dict:  
    print(param_tensor, '\t', state_dict[param_tensor].size())
```

```
State_dict of CNN model:  
conv1.weight      torch.Size([16, 3, 3, 3])  
conv1.bias        torch.Size([16])  
conv2.weight      torch.Size([32, 16, 5, 5])  
conv2.bias        torch.Size([32])  
conv3.weight      torch.Size([128, 32, 3, 3])  
conv3.bias        torch.Size([128])  
fc1.weight        torch.Size([128, 128])  
fc1.bias          torch.Size([128])  
fc2.weight        torch.Size([64, 128])  
fc2.bias          torch.Size([64])  
fc3.weight        torch.Size([10, 64])  
fc3.bias          torch.Size([10])
```

# 加载模型



```
model = CNN()  
#训练模型.....  
  
#保存训练好的参数到文件my_model.pth中  
state_dict = model.state_dict()  
torch.save(state_dict, 'my_model.pth')  
  
#使用保存的模型参数:  
state_dict = torch.load('my_model.pth')  
#把读入的模型参数加载到模型model中:  
model.load_state_dict(state_dict)
```

# 使用预训练模型



```
import torchvision.models as models
resnet18 = models.resnet18()
alexnet = models.alexnet()
vgg16 = models.vgg16()
squeezenet = models.squeezenet1_0()
densenet = models.densenet161()
inception = models.inception_v3()
googlenet = models.googlenet()
shufflenet = models.shufflenet_v2_x1_0()
mobilenet = models.mobilenet_v2()
resnext50_32x4d = models.resnext50_32x4d()
```

# 使用预训练模型



```
resnet50 = models.resnet50(pretrained=True)
```

```
resnet50 = models.resnet50()  
#使用本地磁盘上的模型参数文件  
state_dict = torch.load('resnet50.pth')  
#把读入的模型参数加载到模型model中:  
resnet50.load_state_dict(state_dict)
```

# 使用预训练模型



```
normalize = transforms.Normalize(mean=[0.485, 0.456, 0.406],  
                                std=[0.229, 0.224, 0.225])
```

# 使用ResNet-50作为特征提取器



```
net = models.resnet50(pretrained=False)
state_dict = torch.load('./model/resnet50.pth')
net.load_state_dict(state_dict=state_dict)
for param in net.parameters():
    param.requires_grad=False
```

```
num_classes = 17
featureSize = net.fc.in_features
net.fc = nn.Linear(featureSize, num_classes)
```



# 微调ResNet-50



```
for name, _ in net.named_parameters():  
    print(name)
```

```
layer4.2.conv1.weight  
layer4.2.bn1.weight  
layer4.2.bn1.bias  
layer4.2.conv2.weight  
layer4.2.bn2.weight  
layer4.2.bn2.bias  
layer4.2.conv3.weight  
layer4.2.bn3.weight  
layer4.2.bn3.bias  
fc.weight  
fc.bias
```

# 微调ResNet-50



#这些层不训练

```
exclude_layers = ['layer1', 'layer2', 'layer3']  
for name, param in net.named_parameters():  
    for layer in exclude_layers:  
        if name.startswith(layer):  
            param.requires_grad = False  
            break
```


```
layer4.2.conv1.weight  
layer4.2.bn1.weight  
layer4.2.bn1.bias  
layer4.2.conv2.weight  
layer4.2.bn2.weight  
layer4.2.bn2.bias  
layer4.2.conv3.weight  
layer4.2.bn3.weight  
layer4.2.bn3.bias  
fc.weight  
fc.bias
```

## 3.3 编写复杂的网络

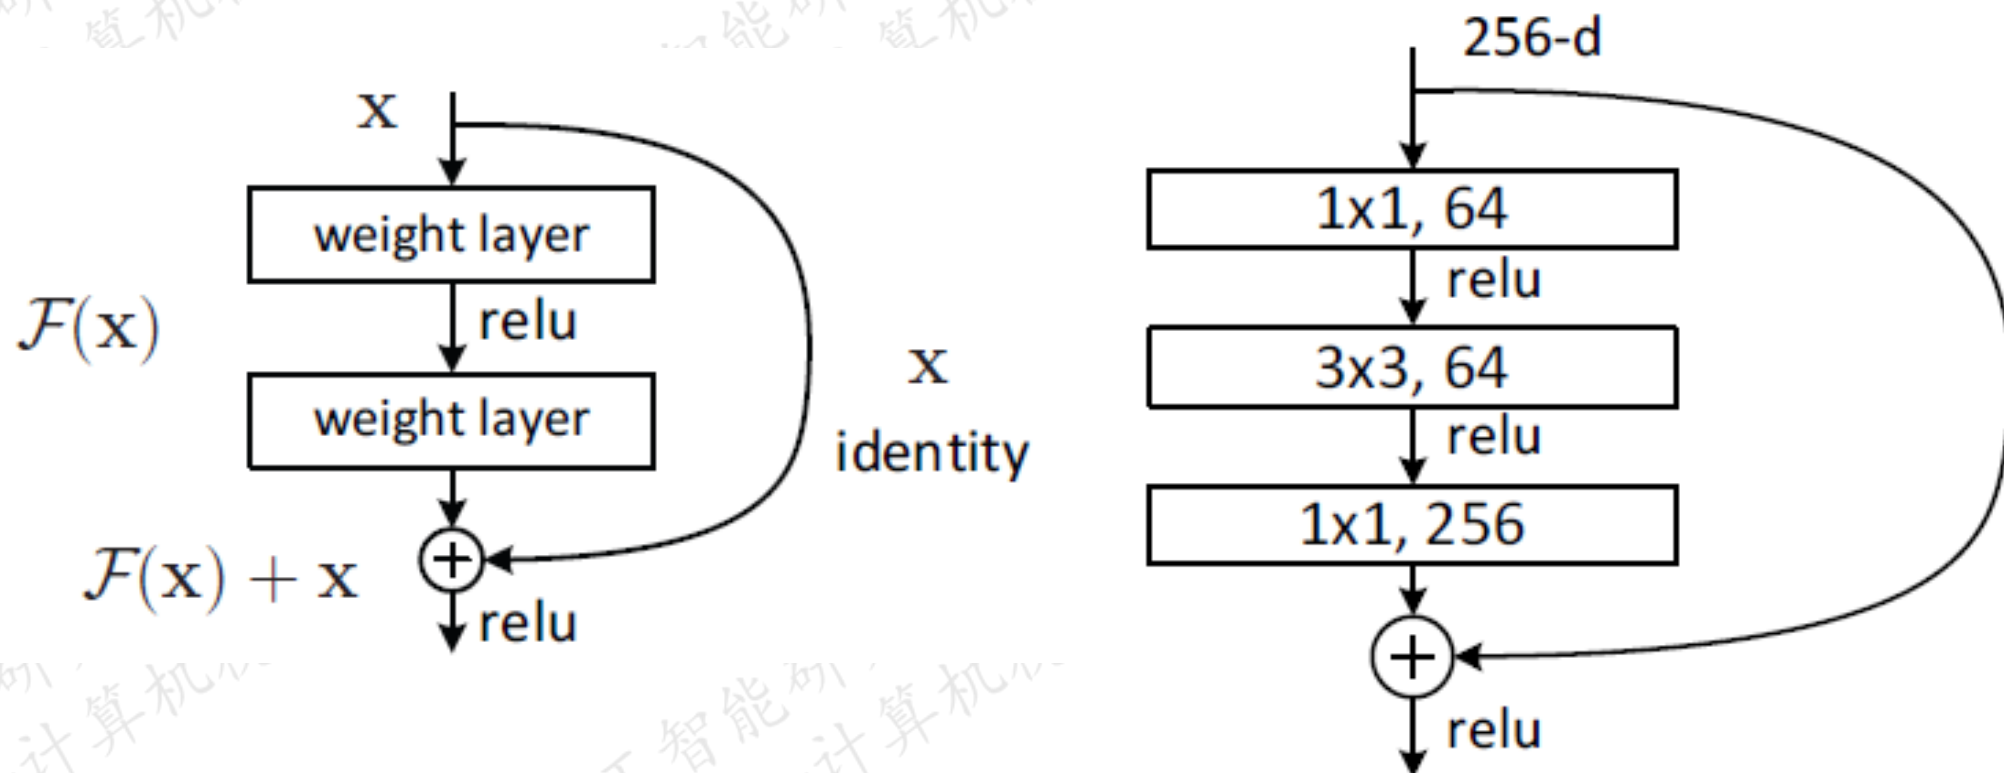
# ResNet



34-layer residual



# Residual Block



# Conv Layer



layer name	output size	18-layer	34-layer	50-layer	101-layer	152-layer
conv1	112×112	7×7, 64, stride 2				
conv2_x	56×56	3×3 max pool, stride 2				
		$\begin{bmatrix} 3 \times 3, 64 \\ 3 \times 3, 64 \end{bmatrix} \times 2$	$\begin{bmatrix} 3 \times 3, 64 \\ 3 \times 3, 64 \end{bmatrix} \times 3$	$\begin{bmatrix} 1 \times 1, 64 \\ 3 \times 3, 64 \\ 1 \times 1, 256 \end{bmatrix} \times 3$	$\begin{bmatrix} 1 \times 1, 64 \\ 3 \times 3, 64 \\ 1 \times 1, 256 \end{bmatrix} \times 3$	$\begin{bmatrix} 1 \times 1, 64 \\ 3 \times 3, 64 \\ 1 \times 1, 256 \end{bmatrix} \times 3$
conv3_x	28×28	$\begin{bmatrix} 3 \times 3, 128 \\ 3 \times 3, 128 \end{bmatrix} \times 2$	$\begin{bmatrix} 3 \times 3, 128 \\ 3 \times 3, 128 \end{bmatrix} \times 4$	$\begin{bmatrix} 1 \times 1, 128 \\ 3 \times 3, 128 \\ 1 \times 1, 512 \end{bmatrix} \times 4$	$\begin{bmatrix} 1 \times 1, 128 \\ 3 \times 3, 128 \\ 1 \times 1, 512 \end{bmatrix} \times 4$	$\begin{bmatrix} 1 \times 1, 128 \\ 3 \times 3, 128 \\ 1 \times 1, 512 \end{bmatrix} \times 8$
conv4_x	14×14	$\begin{bmatrix} 3 \times 3, 256 \\ 3 \times 3, 256 \end{bmatrix} \times 2$	$\begin{bmatrix} 3 \times 3, 256 \\ 3 \times 3, 256 \end{bmatrix} \times 6$	$\begin{bmatrix} 1 \times 1, 256 \\ 3 \times 3, 256 \\ 1 \times 1, 1024 \end{bmatrix} \times 6$	$\begin{bmatrix} 1 \times 1, 256 \\ 3 \times 3, 256 \\ 1 \times 1, 1024 \end{bmatrix} \times 23$	$\begin{bmatrix} 1 \times 1, 256 \\ 3 \times 3, 256 \\ 1 \times 1, 1024 \end{bmatrix} \times 36$
conv5_x	7×7	$\begin{bmatrix} 3 \times 3, 512 \\ 3 \times 3, 512 \end{bmatrix} \times 2$	$\begin{bmatrix} 3 \times 3, 512 \\ 3 \times 3, 512 \end{bmatrix} \times 3$	$\begin{bmatrix} 1 \times 1, 512 \\ 3 \times 3, 512 \\ 1 \times 1, 2048 \end{bmatrix} \times 3$	$\begin{bmatrix} 1 \times 1, 512 \\ 3 \times 3, 512 \\ 1 \times 1, 2048 \end{bmatrix} \times 3$	$\begin{bmatrix} 1 \times 1, 512 \\ 3 \times 3, 512 \\ 1 \times 1, 2048 \end{bmatrix} \times 3$
	1×1	average pool, 1000-d fc, softmax				
FLOPs		$1.8 \times 10^9$	$3.6 \times 10^9$	$3.8 \times 10^9$	$7.6 \times 10^9$	$11.3 \times 10^9$

```
class ResBlock(nn.Module):
    def __init__(self, in_channels, out_channels, downsample = False):
        super(ResBlock, self).__init__()
        stride = 2 if downsample else 1

        self.conv1 = nn.Conv2d(in_channels, out_channels, kernel_size =
3, padding=1, stride = stride, bias = False)
        self.bn1 = nn.BatchNorm2d(num_features = out_channels)

        self.conv2 = nn.Conv2d(out_channels, out_channels, kernel_size =
3, padding=1, stride = 1, bias = False)
        self.bn2 = nn.BatchNorm2d(num_features = out_channels)

        self.relu = nn.ReLU(inplace = True)

        self.downsample = downsample
        if downsample:
            self.sampling_conv = nn.Conv2d(in_channels, out_channels, kernel_size=1,
stride = stride, padding = 0)
            self.sampling_bn = nn.BatchNorm2d(num_features=out_channels)
```

```
def forward(self, x):  
    out = self.conv1(x)  
    out = self.bn1(out)  
    out = self.relu(out)  
  
    out = self.conv2(out)  
    out = self.bn2(out)  
  
    identity = x  
    if self.downsample:  
        identity = self.sampling_conv(identity)  
        identity = self.sampling_bn(identity)  
  
    out += identity  
  
    out = self.relu(out)  
    return out
```



```
class ResLayer(nn.Module):
    def __init__(self, in_channels, out_channels, n_blocks, downsample = True):
        super(ResLayer, self).__init__()
        blocks = [ResBlock(in_channels, out_channels, downsample)]
        for i in range(1, n_blocks):
            blocks.append(ResBlock(out_channels, out_channels))

        self.layer = nn.Sequential(*blocks)

    def forward(self, x):
        return self.layer(x)
```

```

class MyResnet(nn.Module):
    def __init__(self, blocks):
        super(MyResnet, self).__init__()
        channels = 64
        self.conv1 = nn.Conv2d(3, channels, kernel_size = 7, padding = 3, stride = 2,
bias = False)
        self.bn1 = nn.BatchNorm2d(num_features = channels)
        self.pool1 = nn.MaxPool2d(kernel_size=3, stride=2)
        self.gPool = nn.AvgPool2d(kernel_size=7)

        self.layers = [ResLayer(in_channels=channels, out_channels=channels, n_block
s = blocks[0], downsample=False)]

        for b in blocks[1:]:
            self.layers.append(ResLayer(in_channels=channels, out_channels=2*channe
ls, n_blocks = b))
            channels = channels * 2
        self.lin = nn.Linear(in_features=channels, out_features=1000)

```

```
def forward(self,x):  
    x = self.conv1(x)  
    x = self.bn1(x)  
    x = self.pool1(x)  
    for l in self.layers:  
        x = l(x)  
    x = self.gPool(x)  
    x = x.view(-1,x.shape[1])  
    x = self.lin(x)  
  
    return x
```

```
def main():
    resnet = MyResnet([2,2,2,2])

    im = Image.open('Woolsthorpe-Manor.jpg')
    im = im.resize((224,224))
    im = np.array(im, dtype = np.float32) / 255.0
    im_tensor = torch.from_numpy(im)
    im_tensor = im_tensor.permute((2,0,1))
    im_tensor = im_tensor.unsqueeze(dim=0)
    print(im_tensor.size())

    out = resnet(im_tensor)
    print(out.size())

if __name__=='__main__':
    main()
```

```
torch.Size([1,3,224,224])
torch.Size([1,1000])
```

# 练习四



- 1. 复现Deep Residual Learning for Image Recognition论文中的CIFAR10分类识别
- 2. 使用ResNet-50迁移到hymenoptera\_data进行蚂蚁与蜜蜂的识别，输出混淆矩阵，对识别结果进行分析。